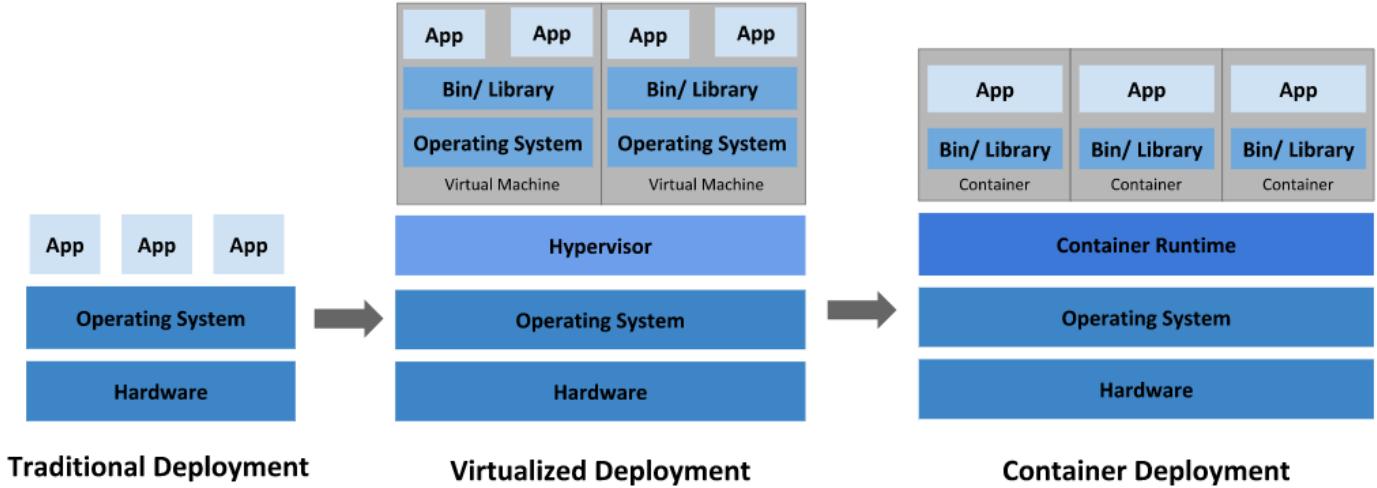


# 6.1) Kubernetes Detailed notes

## Topic:Overview of Containers :-



Why do we need Kubernetes

Question or Concerns about containers

- Containers are a wonderful way of bundling and running your applications, But are they production ready?
- What would happen if the container or the Docker Host goes down?
- How to make containers available 24\*7 ?
- How to handle loads during peak time for the applications ?
- How to replace containers without having downtime with new Docker Image based containers ?
- How to monitor containers?

Solution to above Questions or Concerns

Wouldn't it be good if there is some system which can help for handling all the questions/concerns raised in the above section. That is exactly what Kubernetes does.

- Kubernetes takes care of
  - Scaling requirements
  - failover
  - deployment patterns

- Kubernetes Provides
  - **Service Discovery & Load Balancing**
  - **Storage Orchestration**
  - **Automated rollouts and rollbacks**
  - **Automated bin packing**
  - **Self-Healing**
  - **Secret & Configuration Management**
- Kubernetes is not only for the open source community embraced containers, It is deeply embraced by the Cloud Providers.
  - Amazon Web Service offers **Elastic Kubernetes Services(EKS)**
  - Google Cloud platform offers **Google Kubernetes Engine(GKE)**
  - Microsoft Azure offers **Azure Kubernetes Services(AKS)**

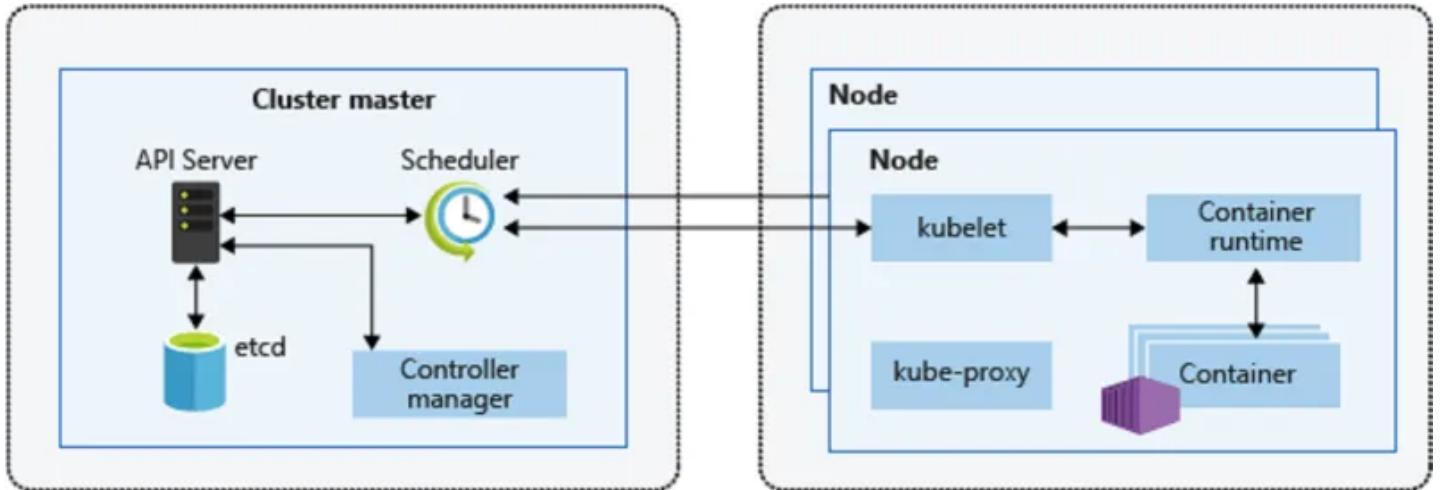
## **Topic:What is Kubernetes:-**

- Kubernetes is a platform that manages container-based applications, their networking and storage components.
- In Kubernetes, we focus on the application workloads rather than the underlying infrastructure.
- Kubernetes provides a declarative approach to deployments, backed by a rich set of APIs for management operations.
- Cluster: Cluster is a collection of compute, storage and networking resources that Kubernetes uses to run workloads.
- Node: It is a single host. Now we can put the cluster as a collection of nodes.

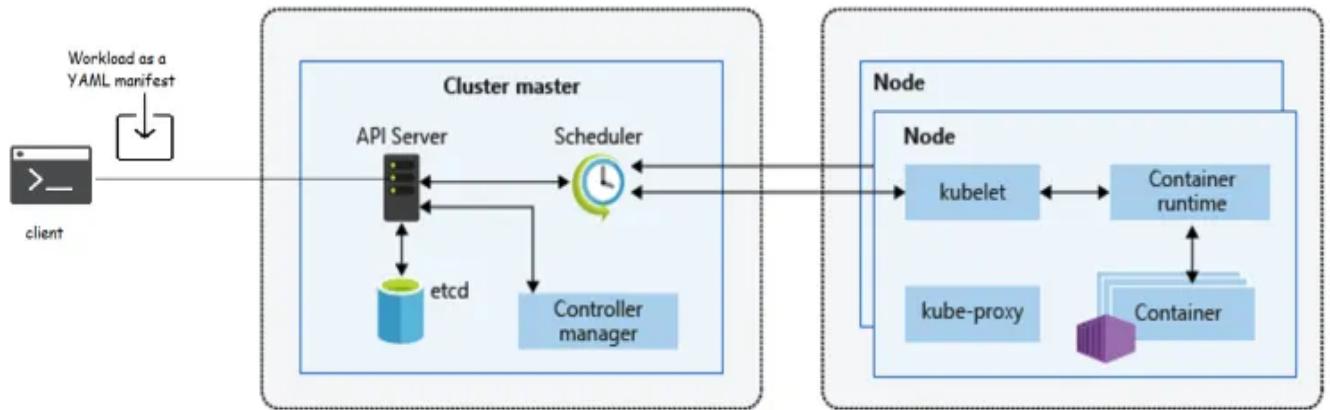
Kubernetes has two kinds of Nodes

- **Master:**
  - Provides core Kubernetes Services and orchestration to application workloads
- **Node:**
  - run your application workloads

Kubernetes Cluster Architecture



What would be our approach with workloads in Kubernetes?



- We would be describing our application as a YAML manifest and pass it to Kubernetes master from some client, and Cluster Master does the rest.



## Pets

### Legacy Infrastructure

Pets are given names like grumpycat.petstore.com

They are unique, lovingly hand raised and cared for  
When they get ill, you nurse them back to health

Infrastructure is a permanent fixture in the data center

Infrastructure takes days to create, are serviced weekly,  
maintained for years, and requires migration projects to move

Infrastructure is modified during maintenance hours and  
generally requires special privileges such as root access

Infrastructure requires several different teams to coordinate and  
provision the full environment

Infrastructure is static, requiring excess capacity to be dormant  
for use during peak periods of demands

Infrastructure is an capital expenditure that charges a fix amount  
regardless of usage patterns



## Cattle

### Cloud-Friendly Infrastructure

Cattle are given numbers like 10200713.cattlerancher.com

They are almost identical to other cattle  
When they get ill, you replace them and get another

Infrastructure is stateless, ephemeral, and transient

Infrastructure is instantiated, modified, destroyed and recreated  
in minutes from scratch using automated scripts

Infrastructure uses version-controlled scripts to modify any  
service without requiring root access or privileged logins

Infrastructure is self-service with the ability to provision  
computing, network and storage services with a single click

Infrastructure is elastic and scales automatically, expanding  
and contracting on-demand to service peak usage periods

Infrastructure is a operating expenditure that charges only for  
services when they are consumed

- k8s is a platform that is used with a huge number of services & capabilities that keeps growing.
- The core functionality of k8s is scheduled workloads in containers across your infra. In addition to this there are many functionalities that k8s brings in
  - Managing storage systems
  - Secret Management
  - Application health & readiness check
  - Replicating application instances
  - Horizontal pod Autoscaling
  - Cluster Autoscaling
  - Naming & Service Discovery
  - Balancing loads
  - Rolling updates
  - Resource Monitoring
  - Log Management
  - Application debugging
  - Authentication & Authorization

# Topic:k8s components:-

<https://kubernetes.io/docs/concepts/overview/components/>

## Kubernetes Master

Cluster:

- It is collection of hosts(nodes) that provide compute, memory, storage and networking resources
- k8s uses these resources to run application workloads.

Master Node:

- The master node is a single host in k8s cluster.
- This can be a physical server or virtual machine
- This is the control plane of k8s. This master is responsible for assigning work (container workloads) to be scheduled on nodes
- Kubernetes master runs on Linux node
- Master Components has i) kube-apiserver ii) etcd iii) kube-scheduler iv) Controller manager v) Cloud-controller-manager

kube-apiserver

- This component is Central to Kubernetes. All communications between all components goes through the kube-apiserver
- This component is frontend of the Kubernetes control plane.
- This component exposes a REST API.
- We would interact with this component using **kubectl** by using the YAML files, which are also referred as manifests

etcd

- etcd stores the entire configuration and the state of the cluster.
- etcd is a consistent and highly available distributed key-value store.

kube-scheduler

- It watches for new work tasks and assigns them to healthy nodes in the cluster

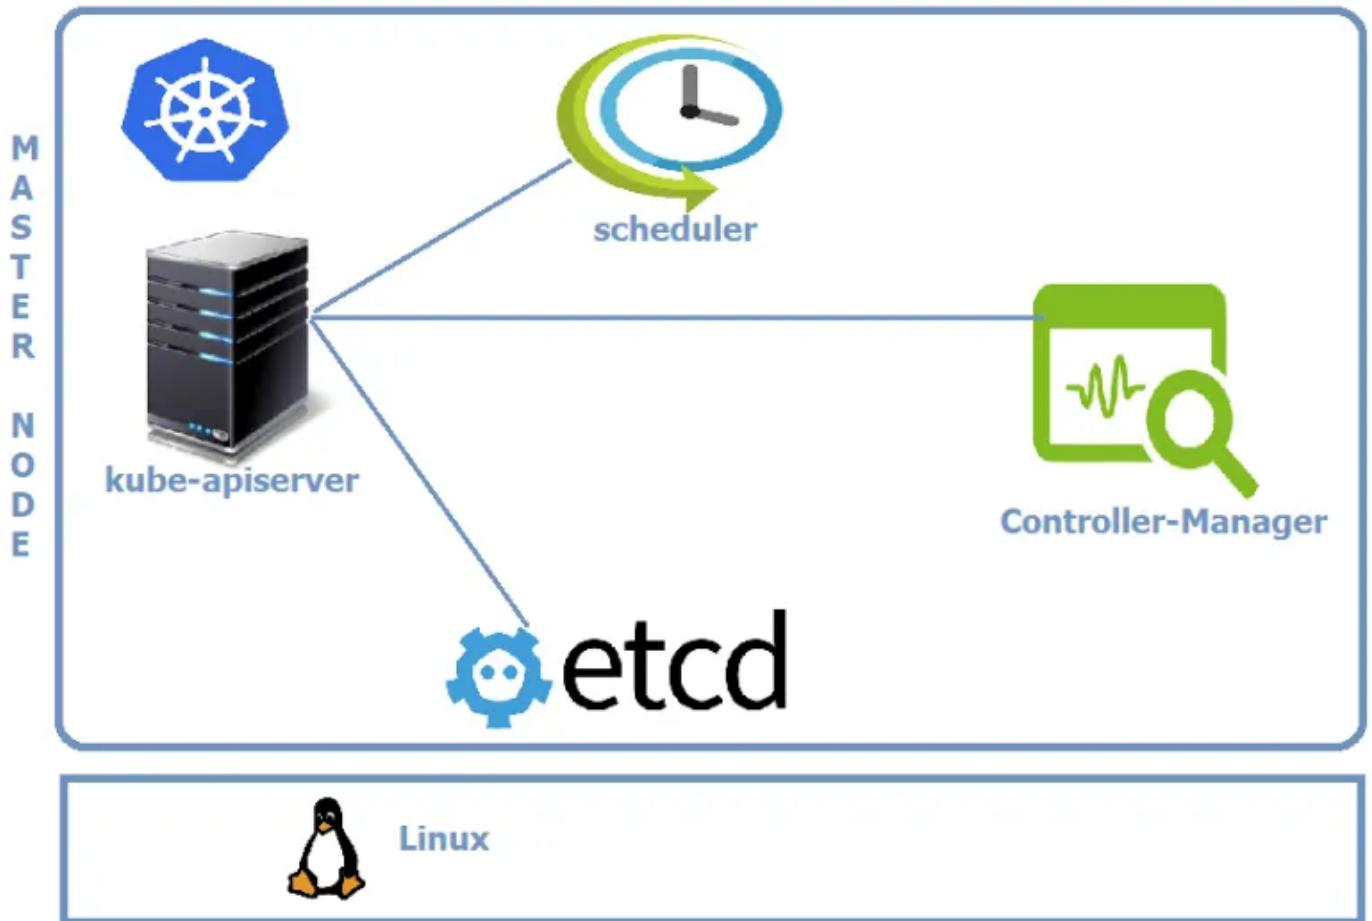
Controller-Manager

- It is responsible for maintaining desired states mentioned in the manifest.
- It looks like single component, but with in it has
  - Node Controller: for noticing & responding when node goes down

- Replication Controller: for maintaining the correct number of pods for every replication controller object.
- Endpoints Controller: Populates the Endpoints object

## Cloud-Controller-manager

- It is responsible for underlying cloud specific controllers.
- If you run the Kubernetes on a supported cloud platform such as AWS, Azure or Google, your control plane runs the Cloud-Controller-Manager.



## Node Components

### Nodes:

- It is a single host in k8s cluster
- This can be physical server or virtual machine

- These are managed by k8s master and they do the work of executing workloads in k8s (They were referred as minions)
- Node is responsible for doing the work assigned.

## kubelet

- This is an agent which runs on each node in the cluster.
- It watches for the instructions from API Server for new work assignments.
- If it can't run the task assigned, it reports back to the master and lets the control plane decide on the actions.
- It is responsible for the node registration process

## Container Runtime

- This is software which is responsible for running containers.
- Some of them are **Docker**, **containerd**, **cri-o**, **rktlet**

## kube-proxy

- Maintains the network rules on nodes
- This is responsible for networking on nodes.

## Cluster DNS

- Every Kubernetes Cluster has an internal DNS service
- This has static IP address that is hardcoded into every Pod on the cluster i.e. all Pods now know how to find the DNS Server
- Services, Stateful Sets and Pods are registered with Cluster DNS.
- Cluster DNS is based on CoreDNS

N  
O  
D  
E



kubelet



container runtime



kube-proxy

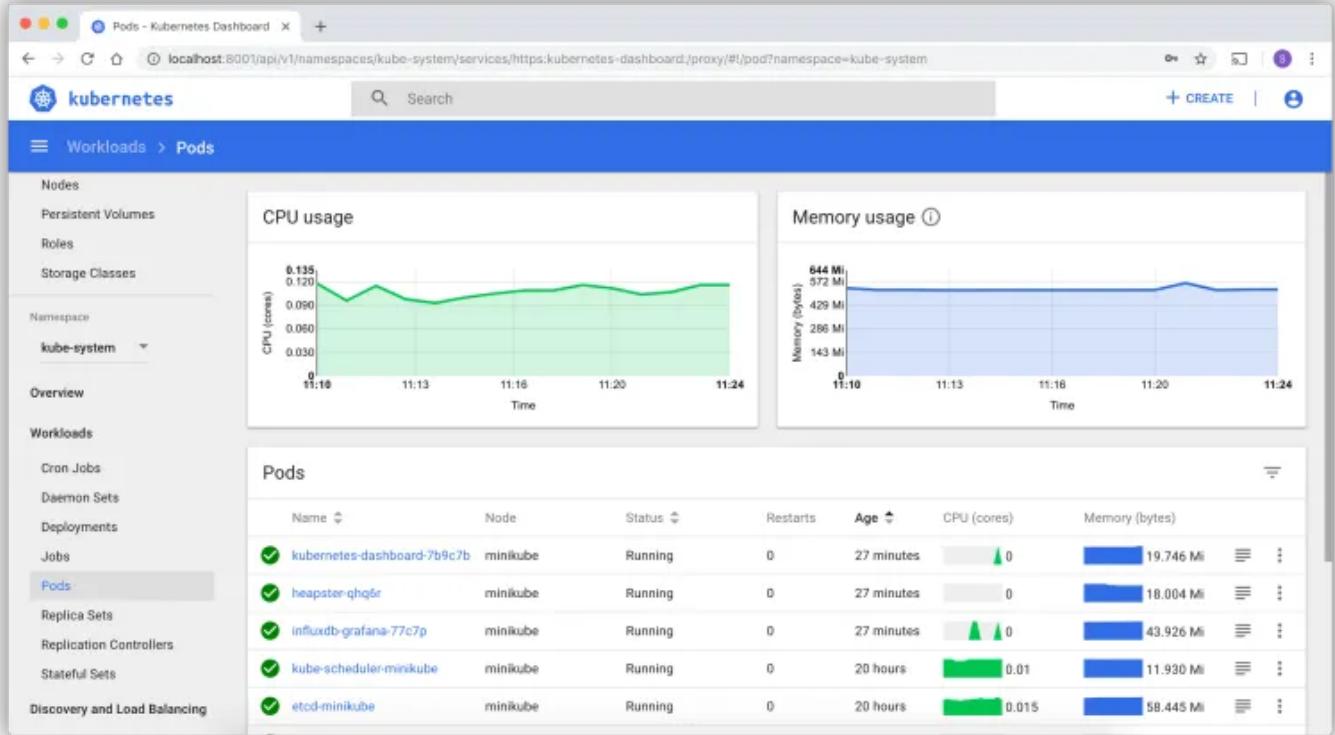


Linux



Interesting ADD ONS

Web UI (Dashboard)



- Web based Kubernetes user interface
- Capabilities:
  - Deployment
  - View cluster resources
  - Troubleshooting

## Container Resource Monitoring

- We will discuss about it in next series of articles

## Cluster Level Logging

- We will discuss about it in next series of articles

## How to Package Apps To Run on Kubernetes

## Deployment

***Scaling, updates and rollbacks***

### Pod

***atomic unit of deployment  
in k8s***

### Container

***Application***

- Package your application as a container image.
  - We would create a Image for the application
  - Push the image to Docker Registry
- Pod will Wrap container
- Pod is deployed into the Kubernetes cluster via Declarative Manifest file.

Pods:-

- A pod is a unit of work in k8s.
- Each pod contains one or more containers.
- Every pod gets a unique ip address and the containers in the pod have the same ip address.

## Topic: K8s Installations:-

- [Refer Here](#) for devops blog installation steps
- [Refer Here](#) for official documentation
- Install docker and kubeadm,kubectl on all the three servers as mentioned in above documentation. Now login into master and execute as [Refereed Over Here](#)
- We need to execute “kubeadm init” it will generate a script,we need to follow according to that and need to create a folder for k8s and configure the network driver for that.
- Now login to kubernetes master and execute “kubeadm init”
- Make node kubeadm init commands login as normal user and create .kube folder and config as below

*Your Kubernetes control-plane has initialized successfully!*

*To start using your cluster, you need to run the following as a regular user:*

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

*Alternatively, if you are the root user, you can run:*

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

*You should now deploy a pod network to the cluster.*

*Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:*

<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

*Then you can join any number of worker nodes by running the following on each as root:*

```
kubeadm join 172.31.41.237:6443 --token 9g3h4c.oxuxjih0ee4f98g6 \  
--discovery-token-ca-cert-hash  
sha256:95646829cf695d5829e9391e1df0453041d797800118f2c6b344e16c3fce041a
```

- As mentioned above if we execute “kubeadm init” it will display like above we need to execute below commands as normal users

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- After that as mentioned in the lines go to the url. This url contain network drivers and that drivers installation commands
  - {<https://kubernetes.io/docs/concepts/cluster-administration/addons/>}
- Generally we have network drivers like Flannel,kniter,Romana..etc in above we have selected weave-network
- I wish to install “weave-network” so we will go to the installation they will give the command
- For install weave-wetwork -network driver we need to execute {kubectl apply -f "[https://cloud.weave.works/k8s/net?k8s-version=\\$\(kubectl version | base64 | tr -d '\n'\)](https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 | tr -d '\n'))"}
- Then join the worker node by following command for ex

```
{kubeadm join 172.31.78.198:6443 --token qbhg9d.v46u4elnoi65bzpp \  
--discovery-token-ca-cert-hash  
sha256:ba79a470521e5f13c4460e7def9a23936890c6bb8a725d86b088d83b334dbd44}
```

- Now join two nodes to cluster by executing kubeadm join command shown above
- After joining nodes, login into master and execute kubectl get nodes

```
ubuntu@ip-172-31-20-190:~$ kubectl get nodes
NAME           STATUS   ROLES      AGE     VERSION
ip-172-31-17-40 Ready    <none>    33s    v1.19.4
ip-172-31-20-190 Ready    master    6m3s    v1.19.4
ip-172-31-27-107 Ready    <none>    79s    v1.19.4
ubuntu@ip-172-31-20-190:~$
```

## Topic: Kubernetes commands:-

- 1) Kubectl version (or) kubectl version -o json {to know the version of k8s}
- 2) kubectl apply -f <yamal file> {To create a POD}
- 3) kubectl get pods {To know number of pods running}

```
root@ip-172-31-17-61:/mahi/k8s/podspecs# kubectl get pods
NAME          READY   STATUS        RESTARTS   AGE
gol-pod       1/1     Running      0          2m41s
hello-world-rest-api-55c678fc85-fc566 0/1     ImagePullBackoff 0          4d22h
root@ip-172-31-17-61:/mahi/k8s/podspecs#
```

- 4) kubectl get pods -o wide {It will display pod clearly}

```
root@ip-172-31-17-61:/mahi/k8s/podspecs# kubectl get pods -o wide
NAME          READY   STATUS        RESTARTS   AGE     IP           NODE   NOMINATED NODE   READINESS GATES
gol-pod       1/1     Running      0          10m    10.44.0.1   ip-172-31-12-116 <none>        <none>
hello-world-rest-api-55c678fc85-fc566 0/1     ImagePullBackoff 0          4d22h  10.47.0.1   ip-172-31-15-137 <none>        <none>
root@ip-172-31-17-61:/mahi/k8s/podspecs#
```

- 5) kubectl describe pod gol-pod {It will describe entire pod details}

```

root@ip-172-31-17-61:/mahi/k8s/podspecs# kubectl describe pod gol-pod
Name:          gol-pod
Namespace:     default
Priority:      0
Node:          ip-172-31-12-116/172.31.12.116
Start Time:    Wed, 07 Apr 2021 09:39:41 +0530
Labels:        <none>
Annotations:   <none>
Status:        Running
IP:           10.44.0.1
 IPs:
  IP:  10.44.0.1
Containers:
  my-gol-cont:
    Container ID:  docker://513a9c89b3e290d5ca3e9c32803add6b6a4fe8c219af9aa9b8baa439dc068b1f
    Image:         qualitythought/gameoflife:07112020
    Image ID:     docker-pullable://qualitythought/gameoflife@sha256:d1e9605cd63903c77ab9cb44a08cb3f5c8da6ec88bbabec2fde4fc0b70b419b7
    Port:          8080/TCP
    Host Port:    0/TCP
    State:        Running
    Started:     Wed, 07 Apr 2021 09:39:52 +0530
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-wltf2 (ro)
Conditions:
  Type      Status
  Initialized  True
  Ready      True
  ContainersReady  True
  PodScheduled  True
Volumes:
  default-token-wltf2:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-wltf2
    Optional:   false
    QoS Class:  BestEffort

```

- 6) kubectl port-forward --address 0.0.0.0 <pod-name> source port(the port where the pod is running):destination port(port where we have to expose the pod)

```

^Croot@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl port-forward --address 0.0.0.0 gameoflife 8080:8080
Forwarding from 0.0.0.0:8080 -> 8080
Handling connection for 8080

```

NOTE:- To know the pod is running on which port etc details execute “kubectl describe <pod-name>” . When ever we do port-forward, default it will try to map on loopback adapter which will work when we are in local machine but whenever you are working on remote machines especially in cloud don’t forget to put ‘0.0.0.0’

- 7) kubectl logs <pod-name> {to view the logs of the pod}

```

root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl logs gameoflife
08-Apr-2021 05:03:22.611 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server version name: Apache Tomcat/9.0.39
08-Apr-2021 05:03:22.616 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server built: Oct 6 2020 14:11:46 UTC
08-Apr-2021 05:03:22.617 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server version number: 9.0.39.0
08-Apr-2021 05:03:22.617 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log OS Name: Linux
08-Apr-2021 05:03:22.617 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log OS Version: 5.4.0-1038-aws
08-Apr-2021 05:03:22.617 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Architecture: amd64
08-Apr-2021 05:03:22.618 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Java Home: /usr/local/openjdk-8/jre
08-Apr-2021 05:03:22.618 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log JVM Version: 1.8.0_272-b10
08-Apr-2021 05:03:22.618 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log JVM Vendor: Oracle Corporation
08-Apr-2021 05:03:22.618 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log CATALINA_BASE: /usr/local/tomcat
08-Apr-2021 05:03:22.619 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log CATALINA_HOME: /usr/local/tomcat
08-Apr-2021 05:03:22.620 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Command Line argument: -Djava.util.logging.config.file=/usr/local/tomcat/conf/logging.properties
08-Apr-2021 05:03:22.620 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Command line argument: -Djava.util.logging.manager=org.apache.juli.ClassLoaderLogManager
08-Apr-2021 05:03:22.620 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Command line argument: -Djdk.tls.ephemeralDHKeySize=2048
08-Apr-2021 05:03:22.621 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Command line argument: -Djava.protocol.handler.pkgs=org.apache.catalina.webresources
08-Apr-2021 05:03:22.621 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Command line argument: -Dorg.apache.catalina.security.SecurityListener.UMASK=0027

```

## 8) kubectl exec <pod-name> -- command {to execute commands in pod using exec }

```

root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl exec gameoflife date
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
Thu Apr  8 07:36:33 UTC 2021
root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl exec gameoflife -- date
Thu Apr  8 07:38:16 UTC 2021
root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl exec gameoflife -- ls
BUILDING.txt
CONTRIBUTING.md
LICENSE
NOTICE
README.md
RELEASE-NOTES
RUNNING.txt
bin
conf
lib
logs
native-jni-lib
temp
webapps
webapps.dist
work
root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# 

```

## 9) kubectl exec -it <pod-name> -- /bin/bash {TO get inside the docker container }

```

root@ip-172-31-17-61:/mahi/k8s/git-repo/KubernetesZone# kubectl exec -it gameoflife -- /bin/bash
root@gameoflife:/usr/local/tomcat# ls
BUILDING.txt CONTRIBUTING.md LICENSE NOTICE README.md RELEASE-NOTES RUNNING.txt bin conf lib logs native-jni-lib temp webapps webapps.dist work
root@gameoflife:/usr/local/tomcat# ll
bash: ll: command not found
root@gameoflife:/usr/local/tomcat# ls -l
total 152
-rw-r--r-- 1 root root 18982 Oct  6 2020 BUILDING.txt
-rw-r--r-- 1 root root 5409 Oct  6 2020 CONTRIBUTING.md
-rw-r--r-- 1 root root 57092 Oct  6 2020 LICENSE
-rw-r--r-- 1 root root 2333 Oct  6 2020 NOTICE
-rw-r--r-- 1 root root 3257 Oct  6 2020 README.md
-rw-r--r-- 1 root root 6898 Oct  6 2020 RELEASE-NOTES
-rw-r--r-- 1 root root 16262 Oct  6 2020 RUNNING.txt
drwxr-xr-x 2 root root 4096 Nov  6 20:27 bin
drwxr-xr-x 1 root root 4096 Apr  8 05:03 conf
drwxr-xr-x 2 root root 4096 Nov  6 20:27 lib
drwxrwxrwx 1 root root 4096 Apr  8 05:03 logs
drwxr-xr-x 2 root root 4096 Nov  6 20:27 native-jni-lib
drwxrwxrwx 2 root root 4096 Nov  6 20:27 temp
drwxr-xr-x 1 root root 4096 Apr  8 05:03 webapps
drwxr-xr-x 7 root root 4096 Oct  6 2020 webapps.dist
drwxrwxrwx 1 root root 4096 Apr  8 05:03 work
root@gameoflife:/usr/local/tomcat# cd webapps
root@gameoflife:/usr/local/tomcat/webapps# ls
gameoflife gameoflife.war

```

## 10) kubectl top nodes {to see the metrics of nodes} && kubectl top pods {to see the metrics of pods}

11) kubectl get rs {to view the replica-set}

NAME	DESIRED	CURRENT	READY	AGE
hello-world-rest-api-55c678fc85	1	1	0	6d8h
replocatio-set	2	2	2	3m39s

12) kubectl describe rs replication-set {to describe the replication set}

kubectl describe rs replication-set					
Name:	replocatio-set				
Namespace:	default				
Selector:	app=gol,version=1.1				
Labels:	<none>				
Annotations:	<none>				
Replicas:	2 current / 2 desired				
Pods Status:	2 Running / 0 Waiting / 0 Succeeded / 0 Failed				
Pod Template:					
Labels:	app=gol				
version=1.1					
Containers:					
gol:					
Image:	qualitythought/gameoflife:07112020				
Port:	8080/TCP				
Host Port:	0/TCP				
Environment:	<none>				
Mounts:	<none>				
Volumes:	<none>				
Events:					
Type	Reason	Age	From	Message	
Normal	SuccessfulCreate	3m59s	replicaset-controller	Created pod: replocatio-set-d2xv6	
Normal	SuccessfulCreate	3m59s	replicaset-controller	Created pod: replocatio-set-ffpfcc	

13) kubectl get rs -o wide {to see rs in detailed}

kubectl get rs -o wide						
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES
gol-rs	2	2	2	80s	gol	qualitythought/gameoflife:07112020

14) kubectl get pods -o wide {to display pods in wide}

kubectl get rs -o wide										
NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES				SELECTOR
gol-rs	2	2	2	80s	gol	qualitythought/gameoflife:07112020				app=gol,version=07112020
kubectl get pods -o wide										
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS	GATES	
gol-rs-4px55	1/1	Running	0	108s	10.47.0.1	ip-172-31-17-40	<none>	<none>	<none>	
gol-rs-mmxf5l	1/1	Running	0	108s	10.44.0.1	ip-172-31-27-107	<none>	<none>	<none>	

15) kubectl get rs -w {to watch the replication set like what is happening}

```
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get rs -w
NAME          DESIRED   CURRENT   READY   AGE
hello-world-rest-api-55c678fc85   1         1         0       6d9h
replocatio-set      2         2         2       17m
replocatio-set      2         2         1       19m
^Croot@ip-172-31-17-61:/mahi/k8s/rs#
```

16) kubectl scale replicaset <rs-name> --replicas=<number>

```
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl scale replicaset replocatio-set --replicas=5
replicaset.apps/replocatio-set scaled
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get ods
error: the server doesn't have a resource type "ods"
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get rs
NAME          DESIRED   CURRENT   READY   AGE
hello-world-rest-api-55c678fc85   1         1         0       6d9h
replocatio-set      5         5         5       30m
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods
NAME                READY   STATUS            RESTARTS   AGE
hello-world-rest-api-55c678fc85-ndd85  0/1    ImagePullBackOff  0          42m
replocatio-set-d2xv6        1/1    Terminating     0          30m
replocatio-set-ffpfcc       1/1    Running          0          30m
replocatio-set-gmdlz        1/1    Running          0          18s
replocatio-set-pdr6q        1/1    Running          0          18s
replocatio-set-sjwrg        1/1    Running          0          18s
replocatio-set-z5cdk        1/1    Running          0          6m10s
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods -o wide
NAME                READY   STATUS            RESTARTS   AGE           IP          NODE   NOMINATED NODE
DINESS GATES
hello-world-rest-api-55c678fc85-ndd85  0/1    ImagePullBackOff  0          42m   10.44.0.1  ip-172-31-12-116  <none>
ne>
replocatio-set-d2xv6        1/1    Terminating     0          30m   10.47.0.1  ip-172-31-15-137  <none>
ne>
replocatio-set-ffpfcc       1/1    Running          0          30m   10.44.0.2  ip-172-31-12-116  <none>
ne>
replocatio-set-gmdlz        1/1    Running          0          26s   10.44.0.4  ip-172-31-12-116  <none>
ne>
replocatio-set-pdr6q        1/1    Running          0          26s   10.44.0.5  ip-172-31-12-116  <none>
ne>
replocatio-set-sjwrg        1/1    Running          0          26s   10.44.0.6  ip-172-31-12-116  <none>
ne>
replocatio-set-z5cdk        1/1    Running          0          6m18s  10.44.0.3  ip-172-31-12-116  <none>
root@ip-172-31-17-61:/mahi/k8s/rs#
```

17) kubectl autoscale rs <name of replicaset> --min=<number> --max=<number>  
--cpu-percent=<percentage> {to configure replicaset for auto scaling using imperative commands}

```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl autoscale rs gol-rs --min=2 --max=5 --cpu-percent=80
horizontalpodautoscaler.autoscaling/gol-rs autoscaled
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME    DESIRED   CURRENT   READY   AGE
gol-rs   2         2         2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME    DESIRED   CURRENT   READY   AGE
gol-rs   2         2         2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get hpa
NAME    REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
gol-rs  ReplicaSet/gol-rs  <unknown>/80%  2          5          2          27s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

18) kubectl get hpa {horizontal pod autoscaler}

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl autoscale rs gol-rs --min=2 --max=5 --cpu-percent=80
horizontalpodautoscaler.autoscaling/gol-rs autoscaled
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
gol-rs     2          2          2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
gol-rs     2          2          2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get hpa
NAME      REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
gol-rs    ReplicaSet/gol-rs <unknown>/80% 2          5          2           27s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

19) kubectl get pods --show-labels {To view names of the labels}

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get pods --show-labels
NAME        READY   STATUS    RESTARTS   AGE   LABELS
gol-rs-label-2hrnf  1/1    Running   0          62s   app=gol,version=07112020
gol-rs-label-9gkjh  1/1    Running   0          62s   app=gol,version=07112020
gol-rs-label-9l9pg  1/1    Running   0          62s   app=gol,version=07112020
gol-rs-label-ln7p7  1/1    Running   0          62s   app=gol,version=07112020
gol-rs-label-spxqj  1/1    Running   0          62s   app=gol,version=07112020
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

20) kubectl get pods --selector="app=<name>" {to view the particular pods with selected name/label}

```

root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods --selector="app=gol"
NAME        READY   STATUS    RESTARTS   AGE
gol-rs-label-2mncp  1/1    Running   0          8m31s
gol-rs-label-dflkp  1/1    Running   0          8m31s
gol-rs-label-gdfj2  1/1    Running   0          8m31s
gol-rs-label-k4rmx  1/1    Running   0          8m31s
gol-rs-label-wrjff  1/1    Running   0          8m31s
root@ip-172-31-17-61:/mahi/k8s/rs#
```

21) Kubectl get pods --selectors in (appname,app-name)

```

ubuntu@ip-172-31-20-190:~/KubernetesZone$ kubectl get pods --selector="app=jenkins"
NAME        READY   STATUS    RESTARTS   AGE
jenkins-rs-label-h854n  1/1    Running   0          82s
jenkins-rs-label-w2x82  1/1    Running   0          82s
ubuntu@ip-172-31-20-190:~/KubernetesZone$ kubectl get pods --selector="app in (jenkins, gol)"
NAME        READY   STATUS    RESTARTS   AGE
gol-rs-label-2hrnf  1/1    Running   0          7m59s
gol-rs-label-9gkjh  1/1    Running   0          7m59s
gol-rs-label-9l9pg  1/1    Running   0          7m59s
gol-rs-label-ln7p7  1/1    Running   0          7m59s
gol-rs-label-spxqj  1/1    Running   0          7m59s
jenkins-rs-label-h854n  1/1    Running   0          108s
jenkins-rs-label-w2x82  1/1    Running   0          108s
ubuntu@ip-172-31-20-190:~/KubernetesZone$
```

22) kubectl api-resources {To know available api-resources}

NAME	SHORTNAMES	APIVERSION	NAMESPACE	KIND
bindings		v1	true	Binding
componentstatuses	cs	v1	false	ComponentStatus
configmaps	cm	v1	true	ConfigMap
endpoints	ep	v1	true	Endpoints
events	ev	v1	true	Event
limitranges	limits	v1	true	LimitRange
namespaces	ns	v1	false	Namespace
nodes	no	v1	false	Node
persistentvolumeclaims	pvc	v1	true	PersistentVolumeClaim
persistentvolumes	pv	v1	false	PersistentVolume
pods	po	v1	true	Pod
podtemplates		v1	true	PodTemplate
replicationcontrollers	rc	v1	true	ReplicationController
resourcequotas	quota	v1	true	ResourceQuota
secrets		v1	true	Secret
serviceaccounts	sa	v1	true	ServiceAccount
services	svc	v1	true	Service
mutatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	MutatingWebhookConfiguration
validatingwebhookconfigurations		admissionregistration.k8s.io/v1	false	ValidatingWebhookConfiguration
customresourcedefinitions	crd,crds	apiextensions.k8s.io/v1	false	CustomResourceDefinition
apiservices		apiregistration.k8s.io/v1	false	APIService
controllerrevisions		apps/v1	true	ControllerRevision
daemonsets	ds	apps/v1	true	DaemonSet
deployments	deploy	apps/v1	true	Deployment
replicasets	rs	apps/v1	true	ReplicaSet
statefulsets	sts	apps/v1	true	StatefulSet
tokenreviews		authentication.k8s.io/v1	false	TokenReview
localsubjectaccessreviews		authorization.k8s.io/v1	true	LocalSubjectAccessReview
selfsubjectaccessreviews		authorization.k8s.io/v1	false	SelfSubjectAccessReview
selfsubjectrulesreviews		authorization.k8s.io/v1	false	SelfSubjectRulesReview
subjectaccessreviews		authorization.k8s.io/v1	false	SubjectAccessReview
horizontalpodautoscalers	hpa	autoscaling/v1	true	HorizontalPodAutoscaler
cronjobs	cj	batch/v1beta1	true	CronJob

23) kubectl get svc {to get service object nodeport, etc}

```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ ls
gol-dev-rs.yaml gol-dev-svc.yaml jenkins-qa-rs.yaml jenkins-test-svc.yaml
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f gol-dev-svc.yaml
service/gol-dev created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
gol-dev   NodePort   10.100.108.213  <none>        8080:32766/TCP  7s
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP      5d23h
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f jenkins-test-svc.yaml
service/jenkins-test created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
gol-dev   NodePort   10.100.108.213  <none>        8080:32766/TCP  2m22s
jenkins-test   NodePort   10.96.163.64   <none>        8080:32765/TCP  11s
kubernetes   ClusterIP   10.96.0.1    <none>        443/TCP      5d23h
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$
```

24) kubectl get namespaces (or) kubectl get ns {To view all name-spaces}

```
root@ip-172-31-27-144:/home/ubuntu# kubectl get ns
NAME          STATUS  AGE
default       Active  6h26m
kube-node-lease  Active  6h26m
kube-public    Active  6h26m
kube-system    Active  6h26m
root@ip-172-31-27-144:/home/ubuntu# |
```

25) kubectl get sts {to get statefuset}

```
root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get sts
NAME          READY   AGE
statefulset-demo 3/3    3m3s
root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
statefulset-demo-0 1/1    Running   0          2m31s
statefulset-demo-1 1/1    Running   0          2m29s
statefulset-demo-2 1/1    Running   0          2m27s
root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE           IP              NODE
STATE NODE  READINESS GATES
statefulset-demo-0 1/1    Running   0          2m45s  192.168.49.114 ip-192-168-34-80.ec2.internal <none>
>      <none>
statefulset-demo-1 1/1    Running   0          2m43s  192.168.21.0 ip-192-168-8-224.ec2.internal <none>
>      <none>
statefulset-demo-2 1/1    Running   0          2m41s  192.168.42.229 ip-192-168-34-80.ec2.internal <none>
>      <none>
root@ip-172-31-27-144:/mahi/k8s/statefull# |
```

26) kubectl get ds {to get demonset}

```
qtkhaja@Azure:~$ kubectl apply -f daemonset.yaml
daemonset.apps/daemonset-demo created
qtkhaja@Azure:~$ kubectl api-resources | grep daemon
daemonsets          ds        apps          true          DaemonSet
qtkhaja@Azure:~$ kubectl get ds
NAME          DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR  AGE
daemonset-demo 3         3        3      3           3          <none>     32s
qtkhaja@Azure:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
daemonset-demo-4vmbn 1/1    Running   0          39s
daemonset-demo-b66cp 1/1    Running   0          39s
daemonset-demo-rwsz9 1/1    Running   0          39s
qtkhaja@Azure:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE           IP              NODE
GATES
daemonset-demo-4vmbn 1/1    Running   0          46s  10.244.1.3  aks-nodepool1-34786526-vmss000001 <none>     <none>
daemonset-demo-b66cp 1/1    Running   0          46s  10.244.2.3  aks-nodepool1-34786526-vmss000002 <none>     <none>
daemonset-demo-rwsz9 1/1    Running   0          46s  10.244.0.17 aks-nodepool1-34786526-vmss000000 <none>     <none>
qtkhaja@Azure:~$ |
```

27) kubectl get deploy {to get deployment set}

```
root@ip-172-31-27-144:/mahi/k8s/controller# kubectl get deploy
NAME          READY   UP-TO-DATE  AVAILABLE  AGE
gameoflife-deploy 3/3    3          3          54m
root@ip-172-31-27-144:/mahi/k8s/controller# |
```

28) Kubectl describe rs <name of rs> {to describe replication set}

```

root@ip-172-31-27-144:/mahik8s/controller# kubectl describe rs gameoflife-deploy
Name:           gameoflife-deploy-754b466d7f
Namespace:      default
Selector:       app=gameoflife,pod-template-hash=754b466d7f,version=1.0
Labels:         app=gameoflife
                pod-template-hash=754b466d7f
                version=1.0
Annotations:    deployment.kubernetes.io/desired-replicas: 3
                deployment.kubernetes.io/max-replicas: 4
                deployment.kubernetes.io/revision: 1
Controlled By: Deployment/gameoflife-deploy
Replicas:       3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 succeeded / 0 Failed
Pod Template:
  Labels:  app=gameoflife
           pod-template-hash=754b466d7f
           version=1.0
  Containers:
    gol-dev:
      Image:      qualitythought/gameoflife:07112020
      Port:       8080/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
Events:
  Type      Reason          Age    From            Message
  ----      ----          ----   ----            -----
Normal   SuccessfulCreate  56m    replicaset-controller  Created pod: gameoflife-deploy-754b466d7f-k7jxd
Normal   SuccessfulCreate  56m    replicaset-controller  Created pod: gameoflife-deploy-754b466d7f-jjsmg
Normal   SuccessfulCreate  56m    replicaset-controller  Created pod: gameoflife-deploy-754b466d7f-9dwpf

```

## Topic:concepts in kubernetes:-

### kubectl

- This is a command-line utility which we would be using to interact with k8s cluster.
- This utility creates api-request internally and forwards them to kube-api server
- [Refer Here](#) for kubectl cheat sheet.

### Kubernetes Objects

- K8s object is anything which can be stored in etcd, As a key-value pair.
- *k8s objects* are persistent entities in the k8s system.
- These objects are used to represent the state of your cluster
- Every k8s object has

- Specification: This is also referred as object Spec
  - Status: This is also referred as object status
- How can we describe a k8s object (Specification):
  - K8s objects can be described in the json format using k8s rest api or in yaml format via kubectl
  - For api version [Refer Here](#)
  - Every k8s object can be described using .yaml (specification) or json when using a rest api and passed across to the kube-api server via kubectl or k8s REST API. The cluster responds back with status
  - While describing object spec, we will have some required fields
    - apiVersion: Which version of k8s api are you using to create this object
    - kind: what kind of k8s object do you want to create
    - metadata: Data which is helpful in uniquely identifying the object, here we will be have *name*, *UID* and optional *namespace*
    - spec: What is your desired state

## K8s Workloads

- Workload is an application running on k8s cluster.
- A Pod has a defined lifecycle. If the pod fails, we need to create the new pod running the same application, K8s has workload resources which can maintain desired state of pods
- Workload resources are
  - Deployments
  - Replica Sets
  - Stateful Sets
  - Daemon Sets
  - Job and Cron Job

## The Pod Manifest (pod yaml file)

- Pods are described in a Pod Manifest which is a text-file representation of k8s API Object.
- Creating a Pod in a imperative kubectl run command

```

ubuntu@ip-172-31-20-190:~$ kubectl run gameoflife --image=qualitythought/gameoflife:07112020
pod/gameoflife created
ubuntu@ip-172-31-20-190:~$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
gameoflife 1/1     Running   0          8s
ubuntu@ip-172-31-20-190:~$ kubectl delete pods gameoflife
pod "gameoflife" deleted
ubuntu@ip-172-31-20-190:~$
```

- Creating a Pod using a manifest file.
- [Refer Here](#) for the pod manifest

```

name: gameoflife
spec:
  containers:
    - image: qualitythought.in/gameoflife:07112020
    - image: qualitythought/gameoflife:07112020
      name: gol
      ports:
        - containerPort: 8080
```

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl apply -f gameoflife-pod.yaml
pod/gameoflife created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
gameoflife 1/1     Running   0          5s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$
```

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl describe pods gameoflife
Name:         gameoflife
Namespace:    default
Priority:    0
Node:        ip-172-31-27-107/172.31.27.107
Start Time:  Sat, 14 Nov 2020 02:23:05 +0000
Labels:      <none>
Annotations: <none>
Status:      Running
IP:          10.44.0.1
IPs:
  IP: 10.44.0.1
Containers:
  gol:
    Container ID:  docker://ba9caf8743cb85713bebafb50de1cd413335786ab2a004392a96b54b24f119af
    Image:        qualitythought/gameoflife:07112020
    Image ID:    docker-pullable://qualitythought/gameoflife@sha256:d1e9605cd63903c77ab9cb44a08cb3f5c8da6ec88bbabe
    c2fde4fc0b70b419b7
    Port:        8080/TCP
    Host Port:  0/TCP
    State:      Running
    Started:    Sat, 14 Nov 2020 02:23:06 +0000
    Ready:      True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-559v4 (ro)
```

- Accessing your pod: {if we can observe the pod ip address and node ip address is different also those are in two different networks. So to access that pod we need to map the ip to port “kubectl port-forward --address 0.0.0.0 gameoflife 8080:8080”}

```
usage:
  kubectl port-forward TYPE/NAME [options] [LOCAL_PORT:]REMOTE_PORT
  [...[LOCAL_PORT_N:]REMOTE_PORT_N]

Use "kubectl options" for a list of global command-line options (applies to all commands).
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl port-forward --address 0.0.0.0 gameoflife 8080:8080
Forwarding from 0.0.0.0:8080 -> 8080
Handling connection for 8080
Handling connection for 8080
Handling connection for 8080
```

← → C ▲ Not secure | 34.218.222.161:8080/gameoflife/

## Welcome to Conway's Game Of Life!

This is a really cool web version of Conway's famous Game Of Life. The Game of Life is a cellular automaton devised by the British mathematician John Horton Conway.

The universe of the Game of Life is an infinite two-dimensional orthogonal grid of square cells, each of which is in one of two possible states, the cells that are directly horizontally, vertically, or diagonally adjacent. At each step in time, the following transitions occur:

- Any live cell with fewer than two live neighbours dies, as if caused by underpopulation.
- Any live cell with more than three live neighbours dies, as if by overcrowding.
- Any live cell with two or three live neighbours lives on to the next generation.
- Any dead cell with exactly three live neighbours becomes a live cell.

New Game

\* Logs information of your pod:

```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl logs gameoflife
14-Nov-2020 02:23:06.767 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server version name: Apache Tomcat/9.0.39
14-Nov-2020 02:23:06.772 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server built: Oct 6 2020 14:11:46 UTC
14-Nov-2020 02:23:06.773 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Server version number: 9.0.39.0
14-Nov-2020 02:23:06.773 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log OS Name: Linux
14-Nov-2020 02:23:06.773 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log OS Version: 5.4.0-1029-aws
14-Nov-2020 02:23:06.774 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Architecture: amd64
14-Nov-2020 02:23:06.774 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log Java Home: /usr/local/openjdk-8/jre
14-Nov-2020 02:23:06.775 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log JVM Version: 1.8.0_272-b10
14-Nov-2020 02:23:06.775 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log JVM Vendor: Oracle Corporation
14-Nov-2020 02:23:06.776 INFO [main] org.apache.catalina.startup.VersionLoggerListener.log CATALINA_HOME: /.../
```

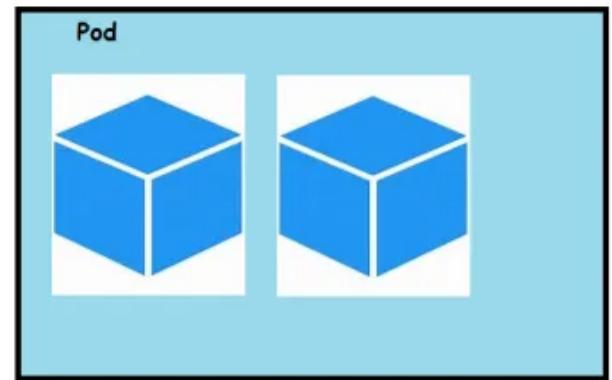
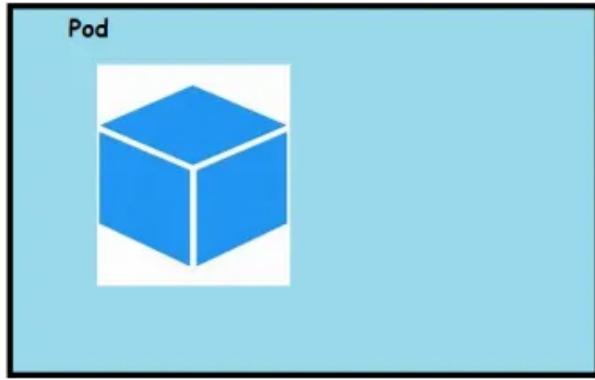
- Running commands in your container with exec

```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl exec gameoflife date
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
Sat Nov 14 02:33:41 UTC 2020
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl exec gameoflife -- date
Sat Nov 14 02:33:57 UTC 2020
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl exec gameoflife -- ls
BUILDING.txt
CONTRIBUTING.md
LICENSE
NOTICE
README.md
RELEASE-NOTES
RUNNING.txt
bin
conf
lib
logs
native-jni-lib
temp
webapps
webapps.dist
work
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$
```

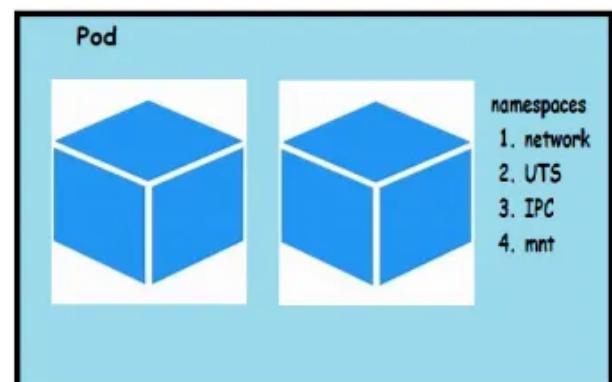
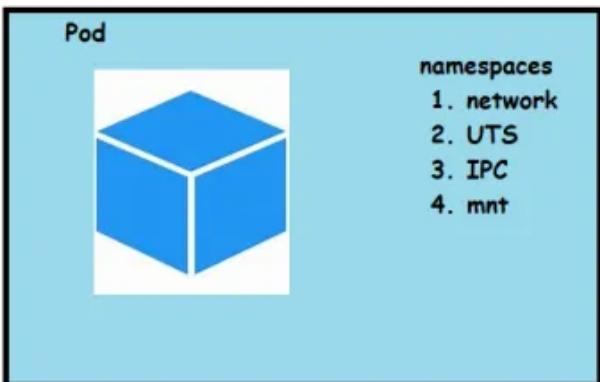
```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/podspecs$ kubectl exec -it gameoflife -- /bin/bash
root@gameoflife:/usr/local/tomcat# ps
bash: ps: command not found
root@gameoflife:/usr/local/tomcat# ls
BUILDING.txt      LICENSE  README.md      RUNNING.txt  conf  logs          temp    webapps.dist
CONTRIBUTING.md  NOTICE   RELEASE-NOTES  bin           lib   native-jni-lib  webapps  work
root@gameoflife:/usr/local/tomcat# cd webapps
root@gameoflife:/usr/local/tomcat/webapps# ls
gameoflife  gameoflife.war
root@gameoflife:/usr/local/tomcat/webapps#
```

## K8s Pods

- Pods are the atomic unit of computing that you create in k8s
- A pod represents an application/microservice running in k8s cluster.
- Pods will have containers in it.



- Each container within a pod runs its own cgroup, but they share a number of Linux namespaces like network, mnt, UTS, IPC, etc.
- These namespaces are for pods. Not for containers.



In k8s to create any object we have two ways:-

Imperative:

- We can create objects by executing commands [Refer Here](#)

Declarative:

- We create an object specification in a yaml file.

- Send/apply this yaml which has a specification & k8s master will do the rest to create/maintain the specification sent
- To create this specification we will be using k8s api reference

## Resource Management:

- Resource Requests for Minimum required resources and maximum limits [Refer Here](#)
- <https://kubernetes.io/docs/concepts/configuration/manage-resources-containers/>
- [Refer Here](#) for resource limits in a pod.

containers:

```
- image: qualitythought/gameoflife:07112020
  name: gol
  resources:
    requests:
      cpu: "500m"
      memory: "128Mi"
    limits:
      cpu: "1000m"
      memory: "512Mi"
  ports:
    - containerPort: 8080
      name: golhttpport
      protocol: TCP
      protocol: TCP
```

- [Refer Here](#) for multi container pod

```
---
apiVersion: v1
kind: Pod
metadata:
  name: gameoflife
spec:
  containers:
    - image: qualitythought/gameoflife:07112020
      name: gol
      resources:
        requests:
          cpu: "500m"
          memory: "128Mi"
        limits:
          cpu: "1000m"
```

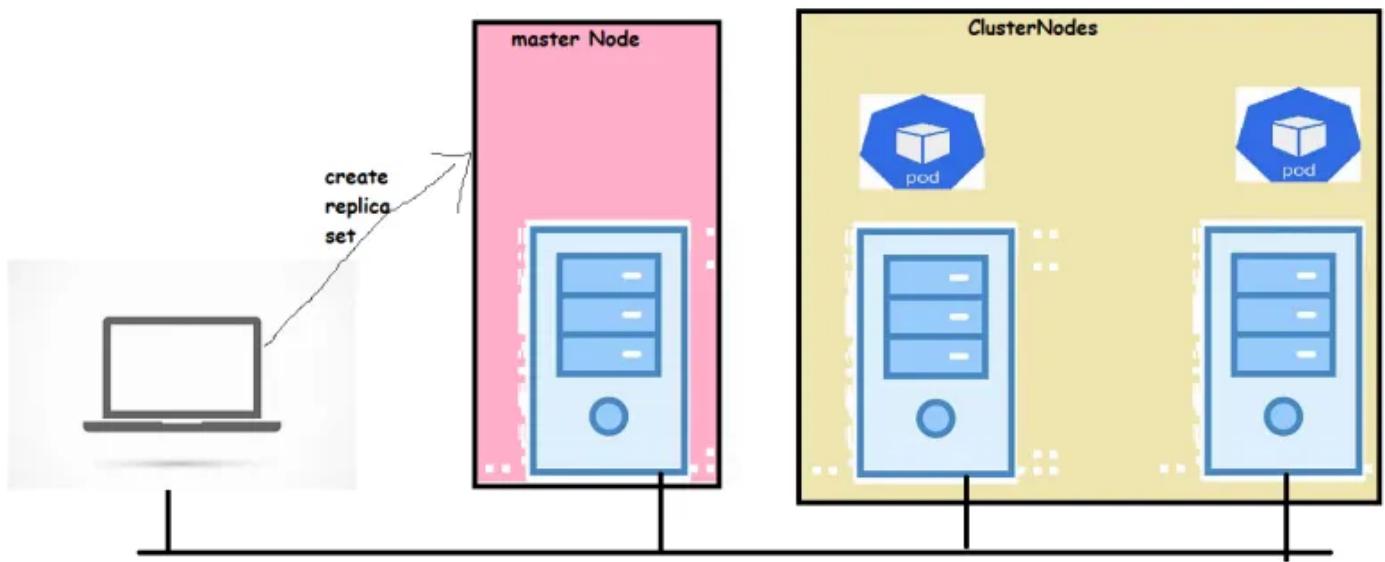
```
    memory: "512Mi"
  ports:
    - containerPort: 8080
      name: golhttpport
      protocol: TCP
    - image: httpd
      name: httpd
      ports:
        - containerPort: 80
          name: httpport
          protocol: TCP
```

## K8s API Groups

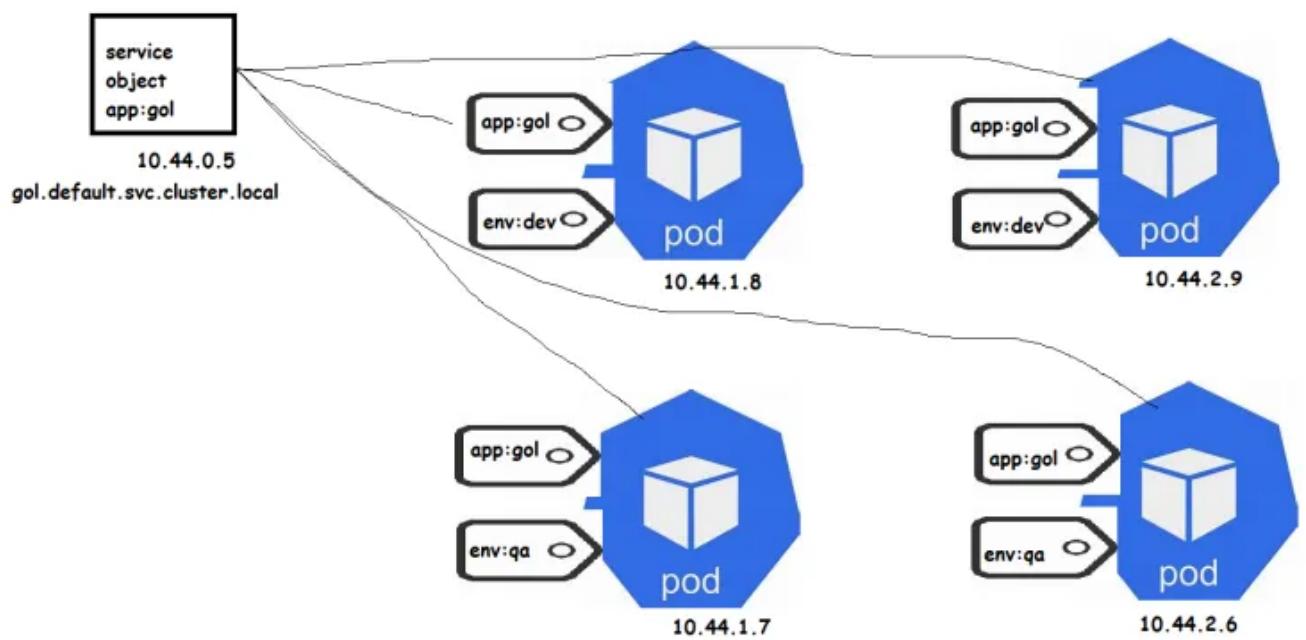
- K8s groups api's as api groups.
- The major purpose is to extend k8s
- [Refer Here](#) for api groups
- To use resources from api groups the naming convention is  
`<GROUP_NAME>/<VERSION>` in yaml file
  - batch/v1
  - extensions/v1beta1
  - rbac.authorization.k8s.io/v1
- To use resource from api groups in the REST api, convention is  
`/apis/<GROUP_NAME>/<VERSION>`
  - /apis/batch/v1
  - /apis/extensions/v1beta1
  - /apis/rbac.authorization.k8s.io/v1
- If the resources are part of core group we can directly write the version `apiVersion:`  
`v1`

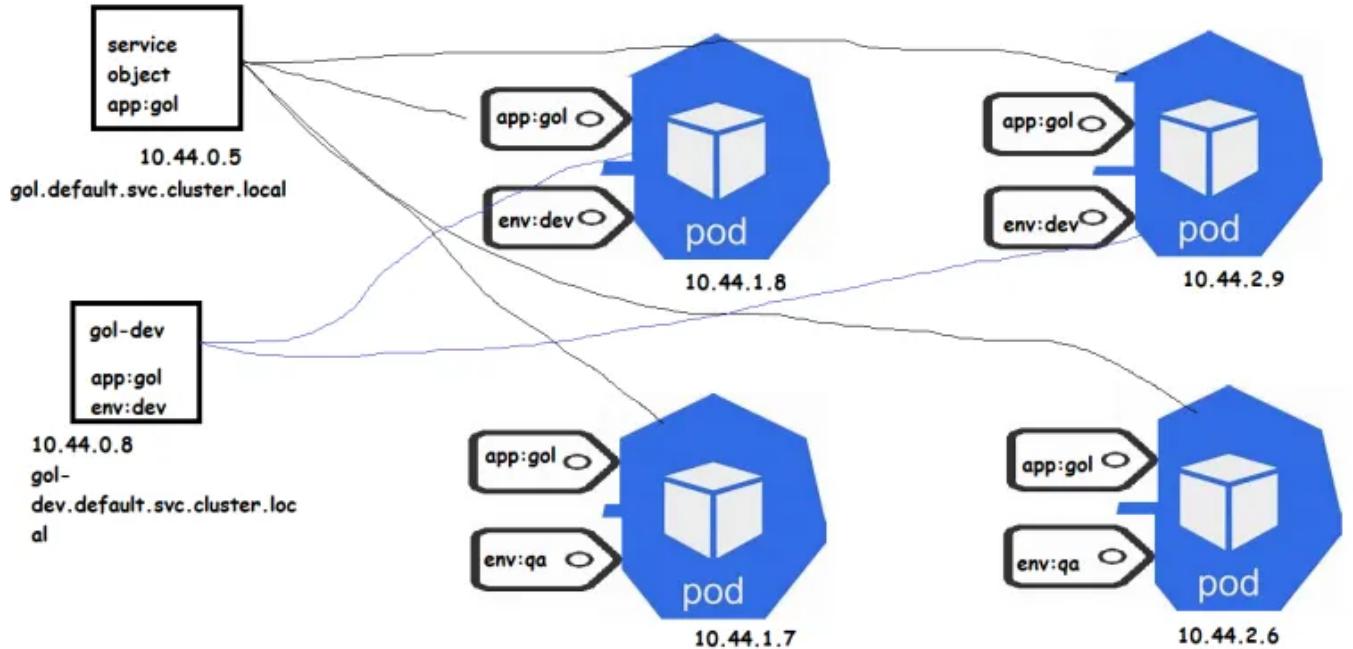
## Topic:- K8s Service

- We have created a replica set with n pods. How do we access the application running inside the pods.
- In normal scenarios for these kinds of problems, we use a load balancer.



- In K8s we use Service Object. A *Service* is a way to create a named label selector. This Service Object will receive an IP Address and also DNS Name





- So now let's create two replica-sets and two services.
- Lets understand service types
  - a. Cluster IP
  - b. Node Port
  - c. load balancer
  - d. external name

EX:-

- now we create 4 objects 2 for replica set in dev environment and test environment. And 2 for services like dev environment and test environment.

```
---
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: gol-dev
spec:
  minReadySeconds: 10
  replicas: 2
  selector:
    matchLabels:
      app: gol
      env: dev
```

```
template:
  metadata:
    labels:
      app: gol
      env: dev
  spec:
    containers:
      - image: qualitythought/gameoflife:07112020
        name: gol-dev
        ports:
          - containerPort: 8080
            protocol: TCP
#####
---
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: jenkins-qa
spec:
  minReadySeconds: 10
  replicas: 2
  selector:
    matchLabels:
      app: jenkins
      env: qa
  template:
    metadata:
      labels:
        app: jenkins
        env: qa
    spec:
      containers:
        - image: jenkins:2.60.3
          name: jenkins-qa
          ports:
            - containerPort: 8080
              protocol: TCP
#####

```

- In the below example in the field of ports.

port belongs to ‘port of your service’  
targetPort belongs to ‘port of your pod’  
nodePort is ‘port of your node. By using which we will access this service’

---

```
apiVersion: v1
kind: Service
metadata:
  name: gol-dev
spec:
  type: NodePort
  selector:
    app: gol
    env: dev
  ports:
    - targetPort: 8080
      port: 8080
      nodePort: 32766
```

#####

---

```
apiVersion: v1
kind: Service
metadata:
  name: jenkins-test
spec:
  selector:
    app: jenkins
    env: qa
  type: NodePort
  ports:
    - targetPort: 8080
      port: 8080
      nodePort: 32765
```

#####

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f jenkins-
jenkins-qa-rs.yaml      jenkins-test-svc.yaml
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f jenkins-qa-rs.yaml
replicaset.apps/jenkins-qa created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f gol-dev-rs.yaml
replicaset.apps/gol-dev created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
gol-dev-pn7xv 1/1     Running   0          6s
gol-dev-wp8fw 1/1     Running   0          6s
jenkins-qa-d77hx 1/1     Running   0          17s
jenkins-qa-zj2fn 1/1     Running   0          17s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get rs
NAME        DESIRED   CURRENT   READY   AGE
gol-dev      2         2         2       12s
jenkins-qa   2         2         2       23s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$
```

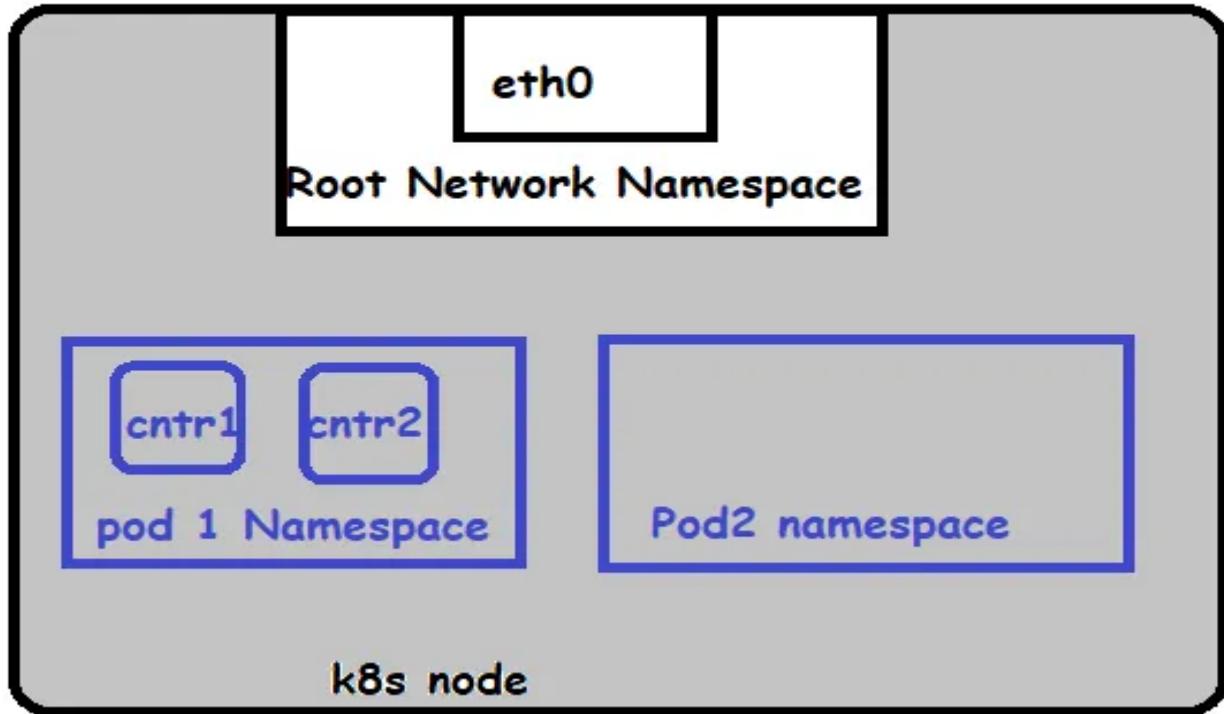
```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ ls
gol-dev-rs.yaml  gol-dev-svc.yaml  jenkins-qa-rs.yaml  jenkins-test-svc.yaml
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f gol-dev-svc.yaml
service/gol-dev created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get svc
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
gol-dev      NodePort    10.100.108.213 <none>           8080:32766/TCP  7s
kubernetes   ClusterIP   10.96.0.1     <none>           443/TCP       5d23h
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl apply -f jenkins-test-svc.yaml
service/jenkins-test created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$ kubectl get svc
NAME        TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
gol-dev      NodePort    10.100.108.213 <none>           8080:32766/TCP  2m22s
jenkins-test  NodePort    10.96.163.64   <none>           8080:32765/TCP  11s
kubernetes   ClusterIP   10.96.0.1     <none>           443/TCP       5d23h
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/service$
```

## Topic: Kubernetes Networking Model (KNM)

- KNM requirements are
  - all pods can communicate with all other pods without using NAT
  - all Nodes can communicate with all Pods without NAT
  - the IP that Pod sees itself is the same IP that other see it as
- Lets understand the following
  - Container to Container networking
  - Pod-to-Pod Networking
  - Pod-to-svc networking
  - Internet to SVC networking

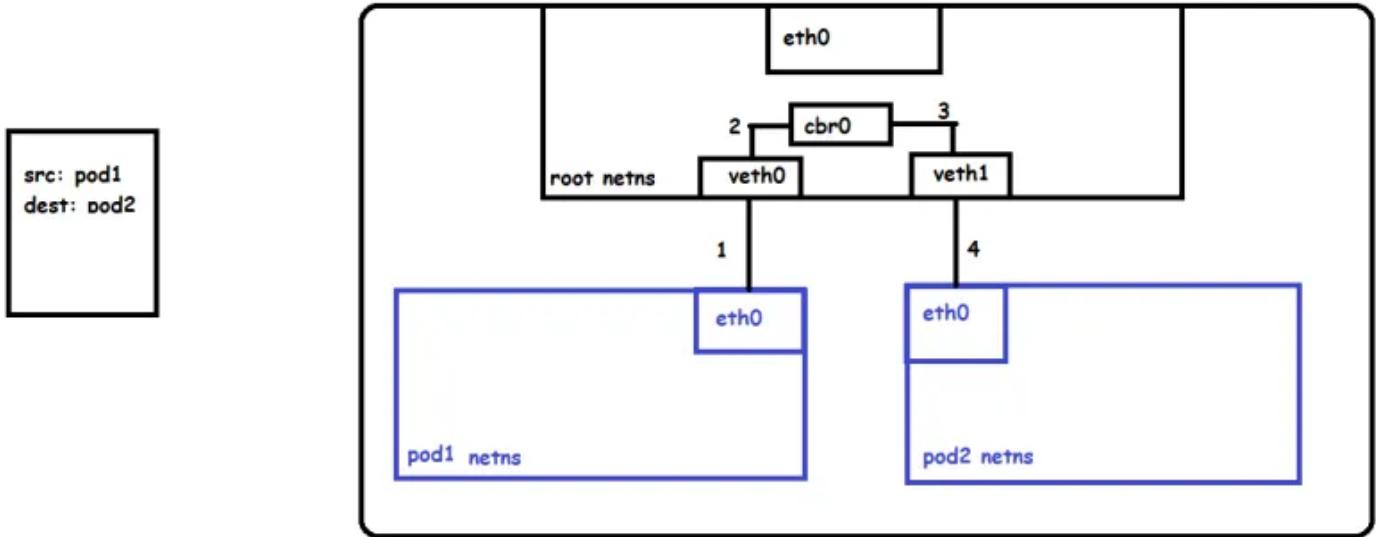
## Container to Container Networking



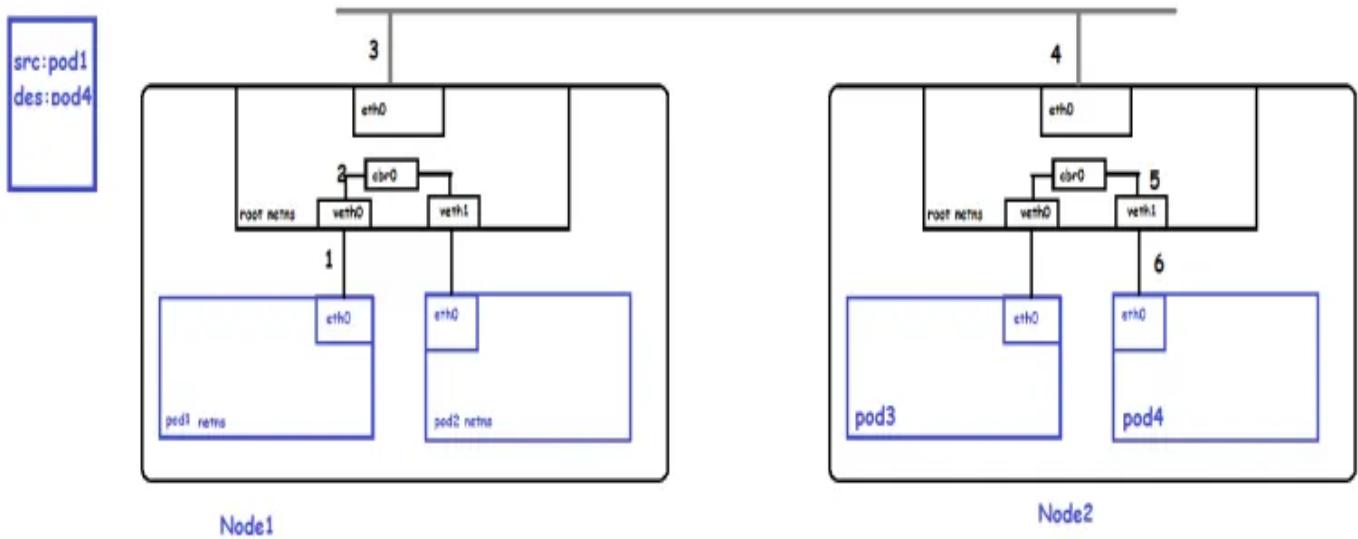
## Pod to Pod Networking

- **Both Pods in Same Node**, will communicate like below one container in pod1 is communicating to another container in pod2. For that node it has a root network, each pod has its own network namespace.
- Because of this network namespace each of this pod will get ethernet interface like ‘eth0,eth0’
- This ‘etc’ is connected to virtual-eather net device ‘veth0’
- These all virtual ethernets which are created in our node n/w name spaces are connected to each other via a bridge ‘cbr0’
- Now we need to send a packet from pod1 to pod2
- Now what will happen is packet will go to eth interface and reached to pod1’s eth0
- Your eth0 is connected to veth0 and it will reach over there and forward to ‘cbr0’
- Then the container bridge will send an ‘arp’ packet to all other networks in that node except source n/w.
- ‘veth1’ will know that this is pod2’s ip address which we are trying to get and it will respond back to container bridge

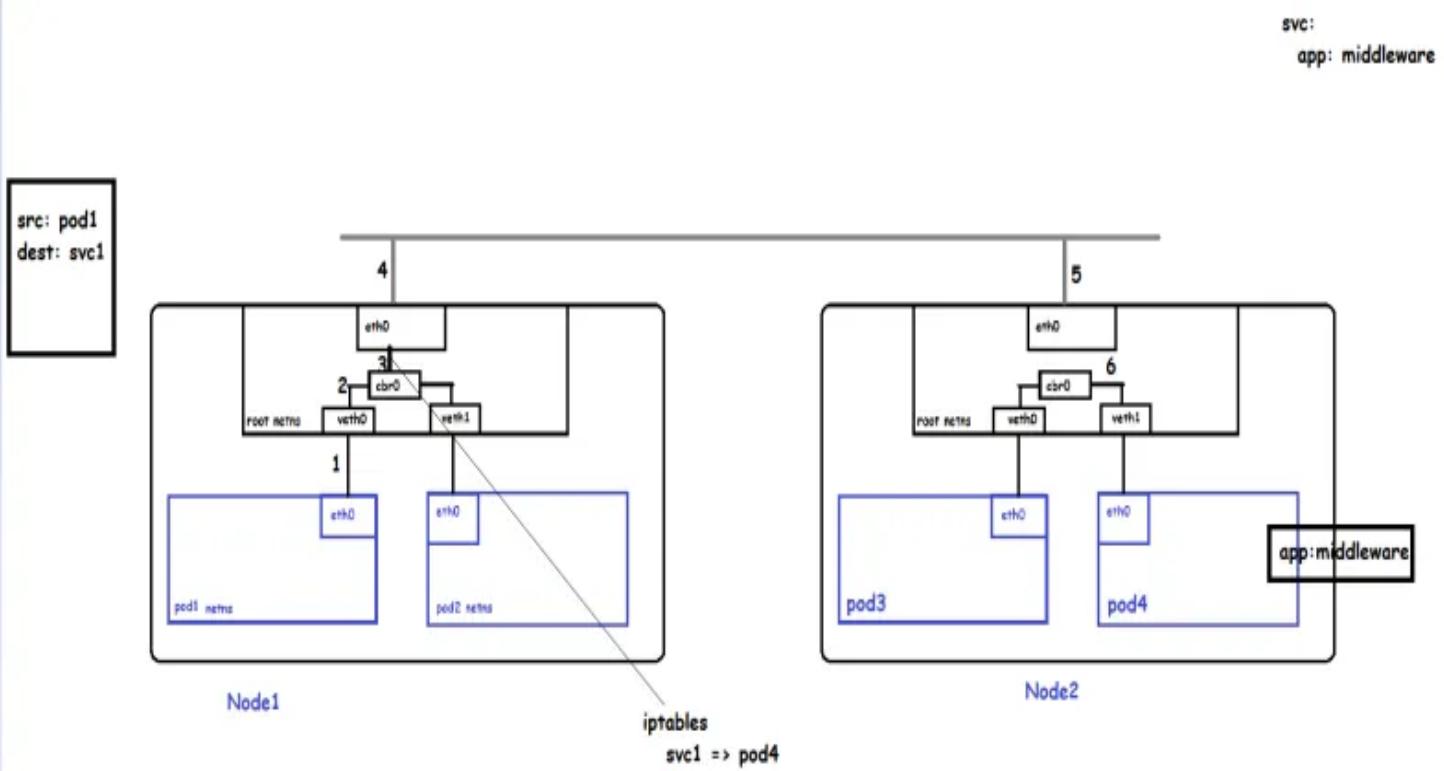
- So this bridge source will send an ‘ARP’ Address resolution protocol to all other networks and any network will respond to that then it will take that and will give to source



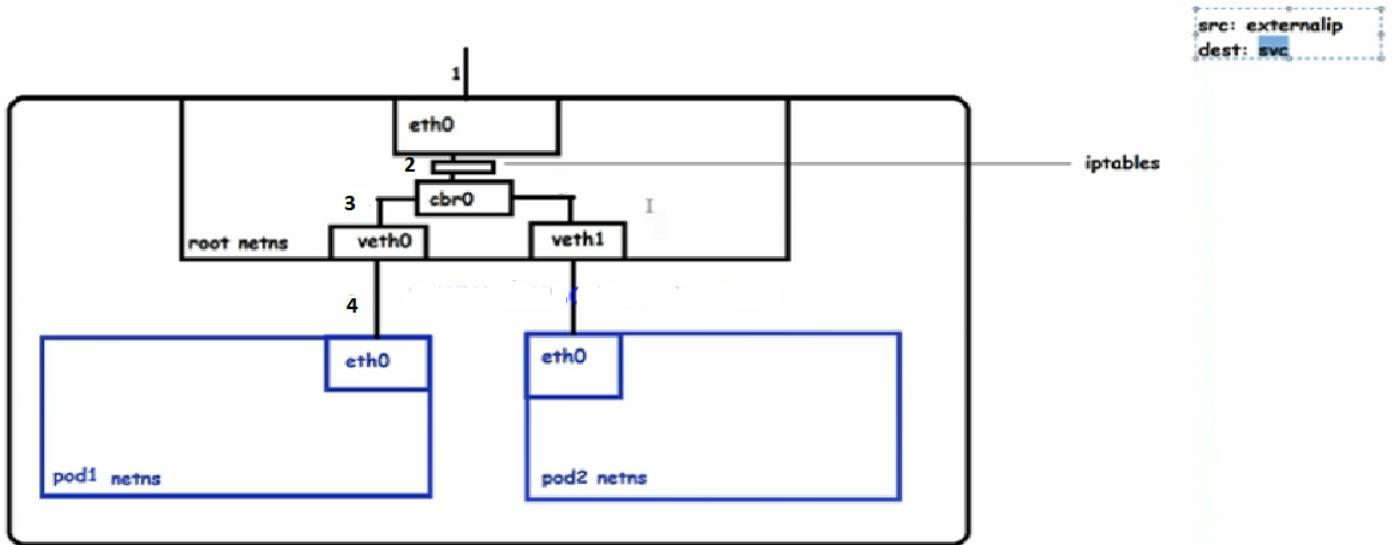
- **Pods on different nodes**, as shown in the below pic, pod1 was trying to communicate with pod4 in node-2.
- For that it will send a packet to eth0 and request will reach to ‘container bridge’ then the bridge will call out for destination ip. If it doesn’t get any response from its own node it will communicate to the external network.
- Basically all nodes are connected via a network. And all the node ethernet will consist of its container's CIDR range. So that it is easy to find out ip of nodes.
- In that pool the matched CIDR node will respond and the packet will reach the destination node. The node-bridge will take that packet to destination pod in this case pod4



- **Pod to Service IP Address**, Here pod1 wants to communicate to service1
- As same as above concepts the packet is reached to eth0 → cb0→ here it will communicate to ‘iptables’ entries
- Because service is an iptable entry that's why service doesn't fail. So here this entry will tell that svc1 means pod4
- Then the packet will reach to destination



- From external world to service, request comes to eth0 then it will check in iptables and as we know all service entries are in iptables and then the request is forwarded to that particular service



## Health Checks of Pod

- When we run our application in k8s, it is kept alive using process health checks.
- However in most cases a simple process might not be enough
  - process has deadlocked and unable to serve requests (process check will return as application is healthy)
- To address this, k8s introduced health checks for application *liveness*
- Liveness Probe:** This probe checks for application specific logic to verify that the application is still running
- Containers that fail liveness probes will be restarted.
- Readiness probe** describes when a container is ready to serve user requests. Containers that fail readiness check are removed from service (load balancers)
- [Refer Here](#) for the liveness probe.

```
---
apiVersion: v1
kind: Pod
metadata:
  name: gameoflife
spec:
  containers:
    - image: qualitythought/gameoflife:07112020
      name: gol
      livenessProbe:
```

```
httpGet:  
  path: /gameoflife  
  port: 8080  
initialDelaySeconds: 5  
timeoutSeconds: 1  
periodSeconds: 10  
failureThreshold: 3  
resources:  
  requests:  
    cpu: "500m"  
    memory: "128Mi"  
  limits:  
    cpu: "1000m"  
    memory: "512Mi"  
ports:  
  - containerPort: 8080  
    name: golhttpport  
    protocol: TCP
```

## Topic:- Activities

### Activity: 1:-

Lets create a pod with the below yaml file.

```
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: gol-pod  
spec:  
  containers:  
    - name: my-gol-cont  
      image: qualitythought/gameoflife:07112020  
      ports:  
        - containerPort: 8080
```

```
ubuntu@ip-172-31-20-190:~/pods$ kubectl apply -f hello-pod.yaml
pod/gol-pod created
ubuntu@ip-172-31-20-190:~/pods$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
gol-pod   1/1     Running   0          16s
ubuntu@ip-172-31-20-190:~/pods$
```

```
ubuntu@ip-172-31-20-190:~/pods$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE      IP           NODE      NOMINATED NODE   READINESS GATES
gol-pod   1/1     Running   0          2m19s   10.44.0.1   ip-172-31-27-107 <none>        <none>
ubuntu@ip-172-31-20-190:~/pods$ kubectl describe pod gol-pod
Name:         gol-pod
Namespace:    default
Priority:    0
Node:        ip-172-31-27-107/172.31.27.107
Start Time:  Fri, 13 Nov 2020 02:56:06 +0000
Labels:      <none>
Annotations: <none>
Status:      Running
IP:          10.44.0.1
IPs:
  IP:  10.44.0.1
Containers:
  my-gol-cont:
    Container ID:  docker://3e2c62abeff5a9858c6b06bad4ad6493d949492ad7f9d896b6e3a510365b8aa5
    Image:        qualitythought/gameoflife:07112020
    Image ID:     docker-pullable://qualitythought/gameoflife@sha256:d1e9605cd63903c77ab9cb44a08cb3f5c8da6ec88bbab
c2fde4fc0b70b419b7
    Port:        8080/TCP
```

```
ubuntu@ip-172-31-20-190:~/pods$ kubectl get pods gol-pod -o yaml
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Pod","metadata":{"annotations":{},"name":"gol-pod","namespace":"default"},"spec":{}}
  containers:[{"image":"qualitythought/gameoflife:07112020","name":"my-gol-cont","ports":[{"containerPort":8080}]}]
  creationTimestamp: "2020-11-13T02:56:05Z"
  managedFields:
  - apiVersion: v1
    fieldsType: FieldsV1
    fieldsV1:
      f:metadata:
        f:annotations:
          .: {}
        f:kubectl.kubernetes.io/last-applied-configuration: {}
    f:spec:
      f:containers:
        k:{"name":"my-gol-cont"}:
          .: {}
        f:image: {}
        f:imagePullPolicy: {}
        f:name: fl
```

```

key: node.kubernetes.io/unreachable
operator: Exists
tolerationSeconds: 300
volumes:
- name: default-token-559v4
  secret:
    defaultMode: 420
    secretName: default-token-559v4
status:
  conditions:
  - lastProbeTime: null
    lastTransitionTime: "2020-11-13T02:56:06Z"
    status: "True"
    type: Initialized
  - lastProbeTime: null
    lastTransitionTime: "2020-11-13T02:56:15Z"
    status: "True"
    type: Ready
  - lastProbeTime: null
    lastTransitionTime: "2020-11-13T02:56:15Z"
    status: "True"
    type: ContainersReady
  - lastProbeTime: null
    lastTransitionTime: "2020-11-13T02:56:05Z"
    status: "True"
    type: PodScheduled
  containerStatuses:
  - containerID: docker://3e2c62abeff5a9858c6b06bad4ad6493d949492ad7f9d896b6e3a510365b8aa5
    image: qualitythought/gameoflife:07112020

```

## Activity:-

## Topic: k8s yaml file

```

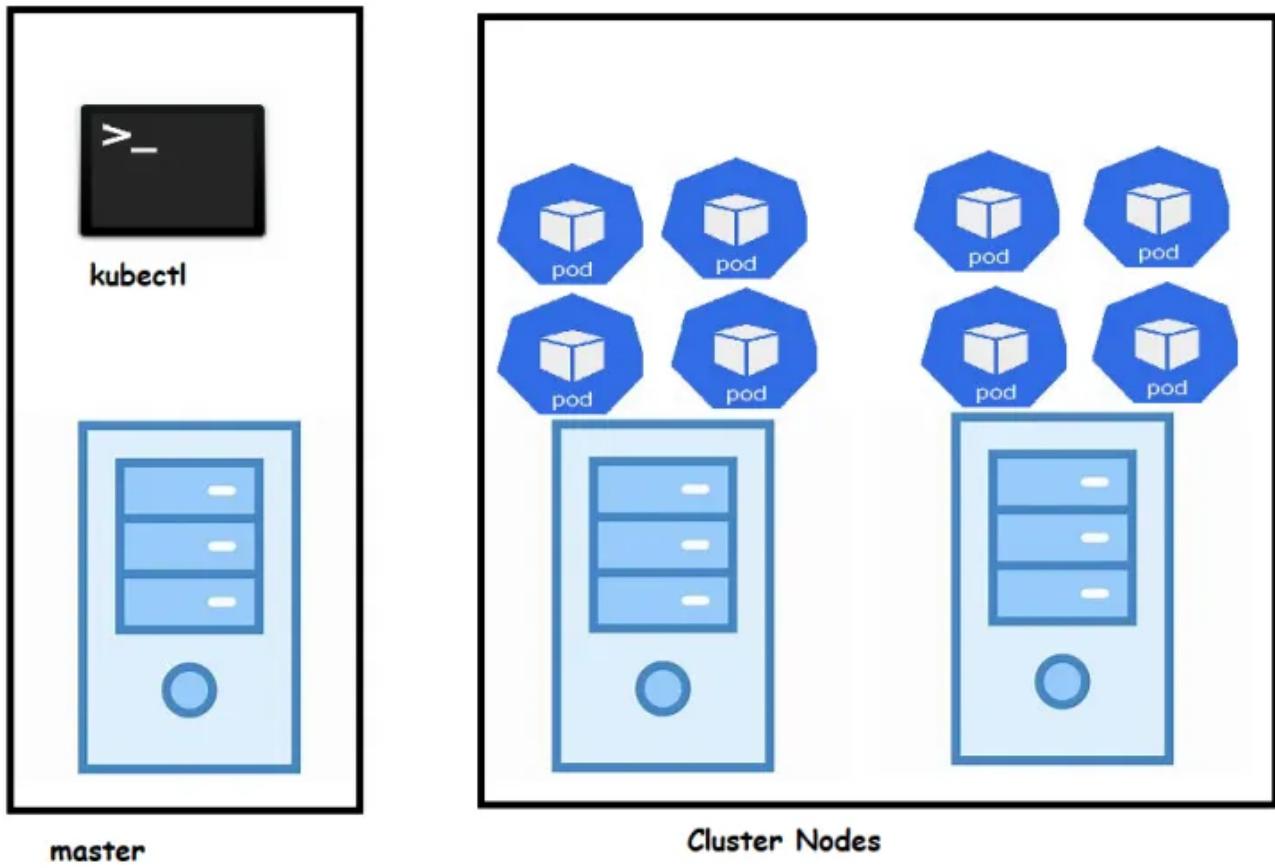
---
apiVersion: v1
kind: Pod
metadata:
  name: gol-pod
spec:
  containers:
  - name: my-gol-cont
    image: qualitythought/gameoflife:07112020
    ports:
    - containerPort: 8080

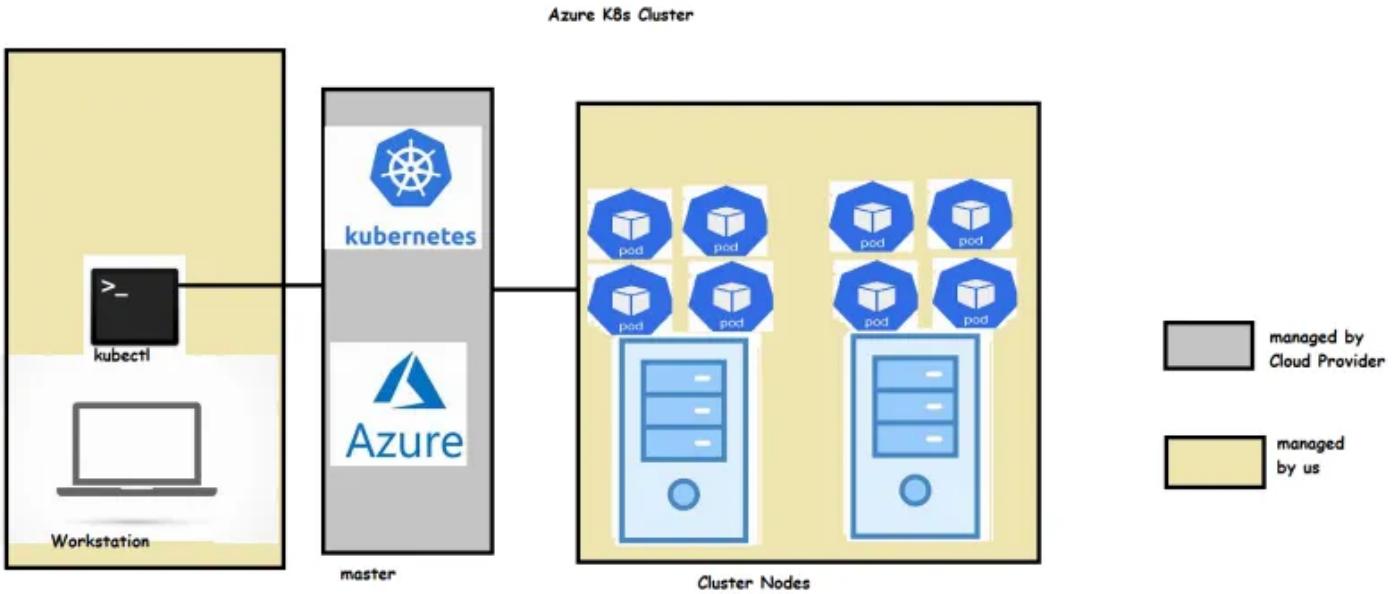
```

# Topic: Azure Kubernetes Services (AKS)

- AKS is a managed k8s cluster in Azure.
- Azure manages master node
- Managed vs UnManage k8s cluster

UnManaged K8s Cluster





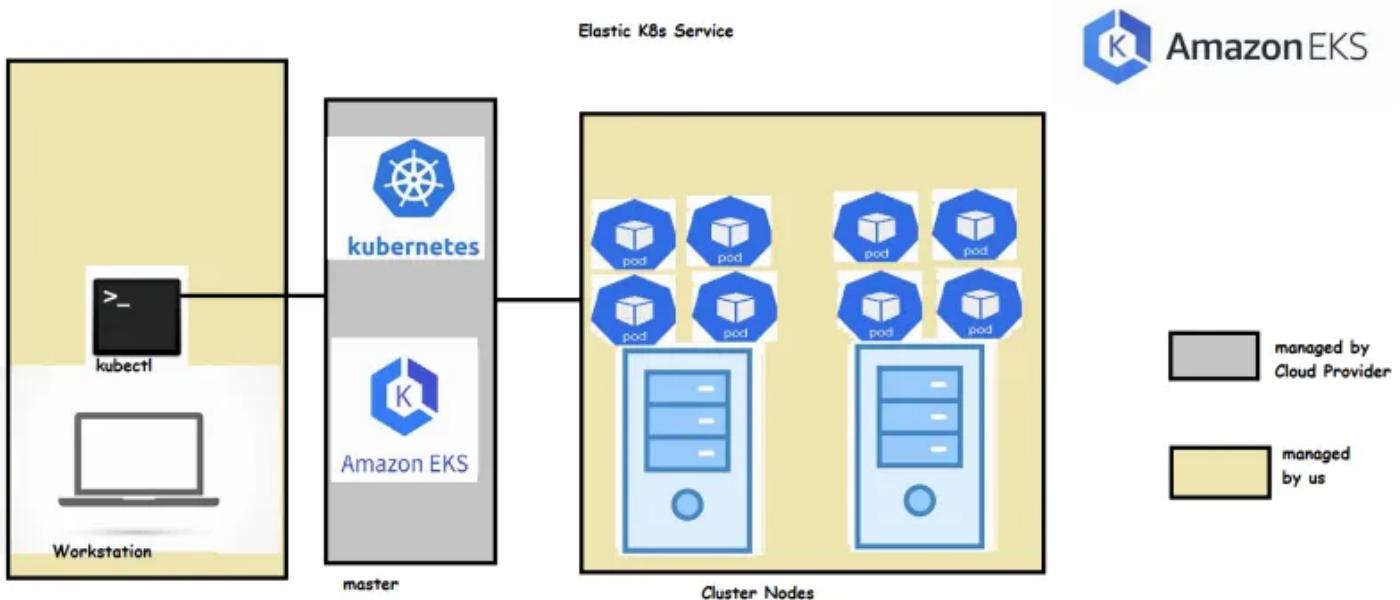
- Lets install AKS Cluster. AKS Cluster can be installed using
  - Portal [Refer Here](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-portal)  
[{https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-portal}](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-portal)
  - CLI [Refer Here](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough)  
[{https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough}](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough)
  - Powershell [Refer Here](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-powershell)  
[{https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-powershell}](https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough-powershell)
- For installing Azure CLI/Powershell [Refer Here](#)

## Topic:Amazon Elastic K8s Service (EKS)

- Lets create a ubuntu ec2 instance
- Install aws cli in the ubuntu
- [Refer Here](#) for setting up eks cluster  
[{https://docs.aws.amazon.com/eks/latest/userguide/getting-started.html}](https://docs.aws.amazon.com/eks/latest/userguide/getting-started.html)
- Install eksctl by following docs [Refer Here](#)  
[{https://docs.aws.amazon.com/eks/latest/userguide/getting-started-eksctl.html}](https://docs.aws.amazon.com/eks/latest/userguide/getting-started-eksctl.html)
- After installing the eks, aws cli and kubectl commands, create k8s cluster by executing the following command

{

```
eksctl create cluster \
--name my-cluster \
--version 1.18 \
--region us-west-2 \
--nodegroup-name linux-nodes \
--nodes 1 \
--nodes-min 1 \
--nodes-max 2 \
--with-oidc \
--ssh-access \
--ssh-public-key dockerclassroom \
--managed
}
```



- 1) Create a vm and install aws cli on that
- 2) Create a user called k8s in IAM and give him admin privileges
- 3) Configure aws cli with the provision of k8s keys
- 4) Install EKS by aws official documentation
- 5) Install Kubectl by using official documents
- 6) Configure managed nodes using docs

## Topic: Namespaces

- K8s is one physical cluster with k8s master and cluster nodes & in this cluster if you want to create multiple virtual clusters (like one for dev, staging, qa...) we can use namespaces.
- Execute kubectl api-resources. The resource with value namespaced

- **true** => will be created per namespace {It means when we create two namespaces that will be different. }
- **false** => will be same/shared across namespaces {if it is false across name-spaces or virtual clusters you can't have different name spaces or nodes. Nodes are same across virtual clusters}

NAME	SHORTNAMES	APIGROUP	NAMESPACED	KIND
bindings			true	Binding
componentstatuses	cs		false	ComponentStatus
configmaps	cm		true	ConfigMap
endpoints	ep		true	Endpoints
events	ev		true	Event
limitranges	limits		true	LimitRange
namespaces	ns		false	Namespace
nodes	no		false	Node
persistentvolumeclaims	pvc		true	PersistentVolumeClaim
persistentvolumes	pv		false	PersistentVolume
pods	po		true	Pod
podtemplates			true	PodTemplate
replicationcontrollers	rc		true	ReplicationController
resourcequotas	quota		true	ResourceQuota
secrets			true	Secret

```
ubuntu@ip-172-31-10-23:~$ kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	8m23s
kube-node-lease	Active	8m24s
kube-public	Active	8m24s
kube-system	Active	8m24s

```
ubuntu@ip-172-31-10-23:~$ kubectl get ns
```

NAME	STATUS	AGE
default	Active	8m31s
kube-node-lease	Active	8m32s
kube-public	Active	8m32s
kube-system	Active	8m32s

```
ubuntu@ip-172-31-10-23:~$
```

- Generally kubectl command will pull the values from namespace *default*. If you want k8s to fetch the information from other namespaces use –namespace flag

```

ubuntu@ip-172-31-10-23:~$ kubectl get pods
No resources found in default namespace.
ubuntu@ip-172-31-10-23:~$ kubectl get pods --namespace kube-system
NAME                  READY   STATUS    RESTARTS   AGE
aws-node-mndqp       1/1     Running   0          3m35s
coredns-559b5db75d-cq59l 1/1     Running   0          10m
coredns-559b5db75d-djrmv 1/1     Running   0          10m
kube-proxy-nckts     1/1     Running   0          3m35s
ubuntu@ip-172-31-10-23:~$
```

NOTE:- If we need to k8s will pull the values from some particular name-space, Then you can set something called as context (context is nothing but some values stored in kubectl config file `~/.kube/config`)

```

apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: LS0tLS1CRUDJTiBDRVJUSUZJQ0FURSOtLS0tCK1JSUN5RENQWJD0F3SUJBZ01CQURBTkJna3Foa21HOXcwQkFRc0ZBREFWTVJNdo
VRWURWUVFERXdwmRXsmwKY201bGRHVnpNojRYFRJRe1EUxhNakV3TwpVek5sb1hEVE14TURRe1ERXdnalV6tmxvd0ZURVRN0kVHOTFVRQpBe1LYTNWa/pY5nVaWFJsV3yPd
Q0FTSXdeUVIKs29aswh2Y05BUUVCQ1FBRGdnRVBBRENDQVFvQ2dnuRJBT0pYCrnJpZw9nQlo2bwg4U2RmwVvYnKxx1b21ta0xeNm44eWv1s1VKcvB0WEUyOGJTB21Qe1d1mjYxSE
hFL2RwbVV0dFIKVnNjwG5TQ1zuvXM1Unp1Yml0WvZsNm9HY1pqcZrja0xJUEZQcTBWQzFjZVvKQ0Y3F4ZTjzQ1FsVAo0d1JGNFFmU0svQ1RoQUhLNwZrVkdTwU9u
cg9nVG9sMgtjTU1ILz1rakFxWpvt1jkswsyMGNDUFV4bdVsieg5Xcmkzdm1hv1IPZUNUbFzob29mRtk3Tkdnu0fMUwLwZ3RxWnMwd2xZdkpuNmlrZwdFZwpjbVvZdxJzv3BDQk
MOV1YKbDuySFdzOHZ5WEZDQjIxQVRRXF0QzJrau4M1JmK3RBuWZFS1NEZ0ppRkpINEFd1R2wxNyYXNENwtZqxu3dQpTb1ZIVHzuZzJpc3A5SxUyTHMwQ0f3RUFByU1qTUNF
d0RnWURWU1BQQVFILOJBURBZ0trTUE4R0ExVWRFd0VCC193UzNQ1CQWY4d0RRWUpLb1pJaHzjTKFRRUXCUUFEZ2dFQkFOckdISU9sNz1jajdMb1hc11twmRXSkk2ZH1KNW
V1c0Y4awk0Vudjv1dmv1pvQ1M3Vkp3RktveUNQZFF6vUZnbHhZckppRHhsVnVwWhxMTZubX1NTFY2amViwgpwQThHSkg1l1dzVzdqcnMls3h20T11cFRRUUxrQTZBaHNYTGY
VTY4OFF6NERzam5rWfpjeGJGUmfLMX13ekc5CnaxbCTWu0xdGdsZk5iRVgwSm5Hb2tvcnVybHBCZVfyc0w5R3FnS0FUZ1lnSnhtdkVxb05aQmdwSFBiYTFNbWCKSEgzsWu1K0
RIRG80aw1RV2tneEpRY1NDVkvRq1JuVkh1L2t50mRTcd1ju9cdFRTwMzpQzV1RysvcGdTWTduaAp2enc0MER0VklYamEveHlsTT1MeGradDFHZFBpc1EzRWE2djVUVUFKngz
Yk11TkxGL1hsdkN3Z0jJUT0KLS0tLS1F0qQ0VS17GSUNBVETULs0tLQo=
  server: https://039758468CAD1FCDA01552014E982FF5.gr7.us-east-1.eks.amazonaws.com
  name: mahi-k8s.us-east-1.eksctl.io
contexts:
- context:
  cluster: mahi-k8s.us-east-1.eksctl.io
  user: k8s@mahi-k8s.us-east-1.eksctl.io
  name: k8s@mahi-k8s.us-east-1.eksctl.io
current-context: k8s@mahi-k8s.us-east-1.eksctl.io
kind: Config
preferences: {}
users:
- name: k8s@mahi-k8s.us-east-1.eksctl.io
user:
  exec:
    apiVersion: client.authentication.k8s.io/v1alpha1
    args:
      - eks
      - get-token
      - --cluster-name
      - mahi-k8s
      - --region
      - us-east-1
```

## Upgrading

- Upgrading to published versions is easier in the cloud.
  - In AKS, we can upgrade the nodepools and k8s seperately
  - In the case of eksctl

```

ubuntu@ip-172-31-10-23:~$ eksctl upgrade --help
Upgrade resource(s)

Usage: eksctl upgrade [flags]

Commands:
  eksctl upgrade cluster           Upgrade control plane to the next version
  eksctl upgrade nodegroup         Upgrade nodegroup

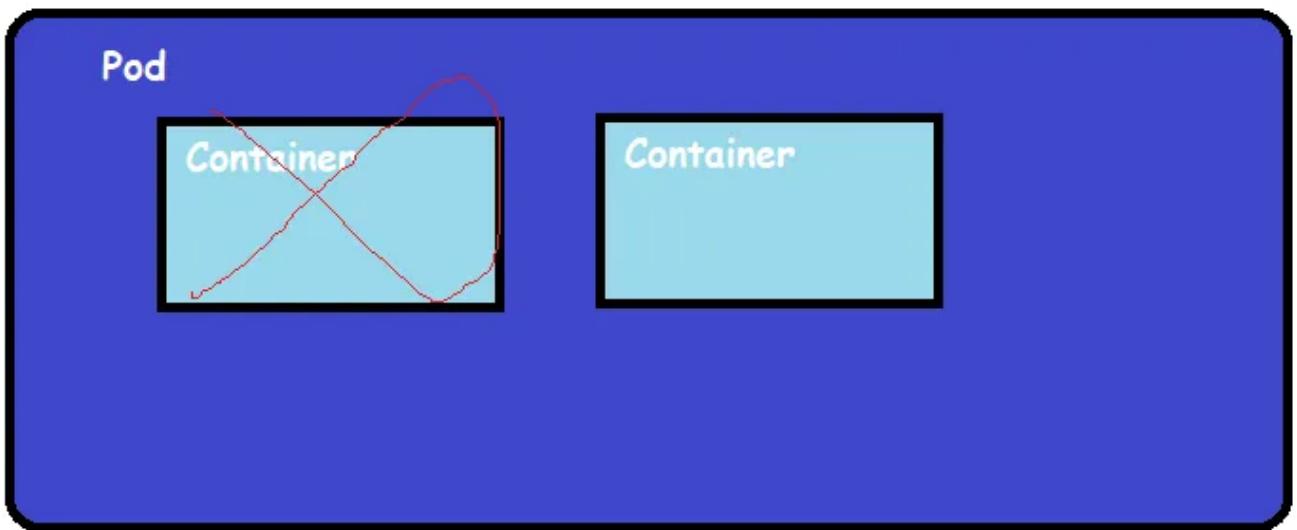
Common flags:
  -C, --color string   toggle colorized logs (valid options: true, false, fabulous) (default "true")
  -h, --help            help for this command
  -v, --verbose int    set log level, use 0 to silence, 4 for debugging and 5 for debugging with AWS debug logging (default 3)

Use 'eksctl upgrade [command] --help' for more information about a command.

```

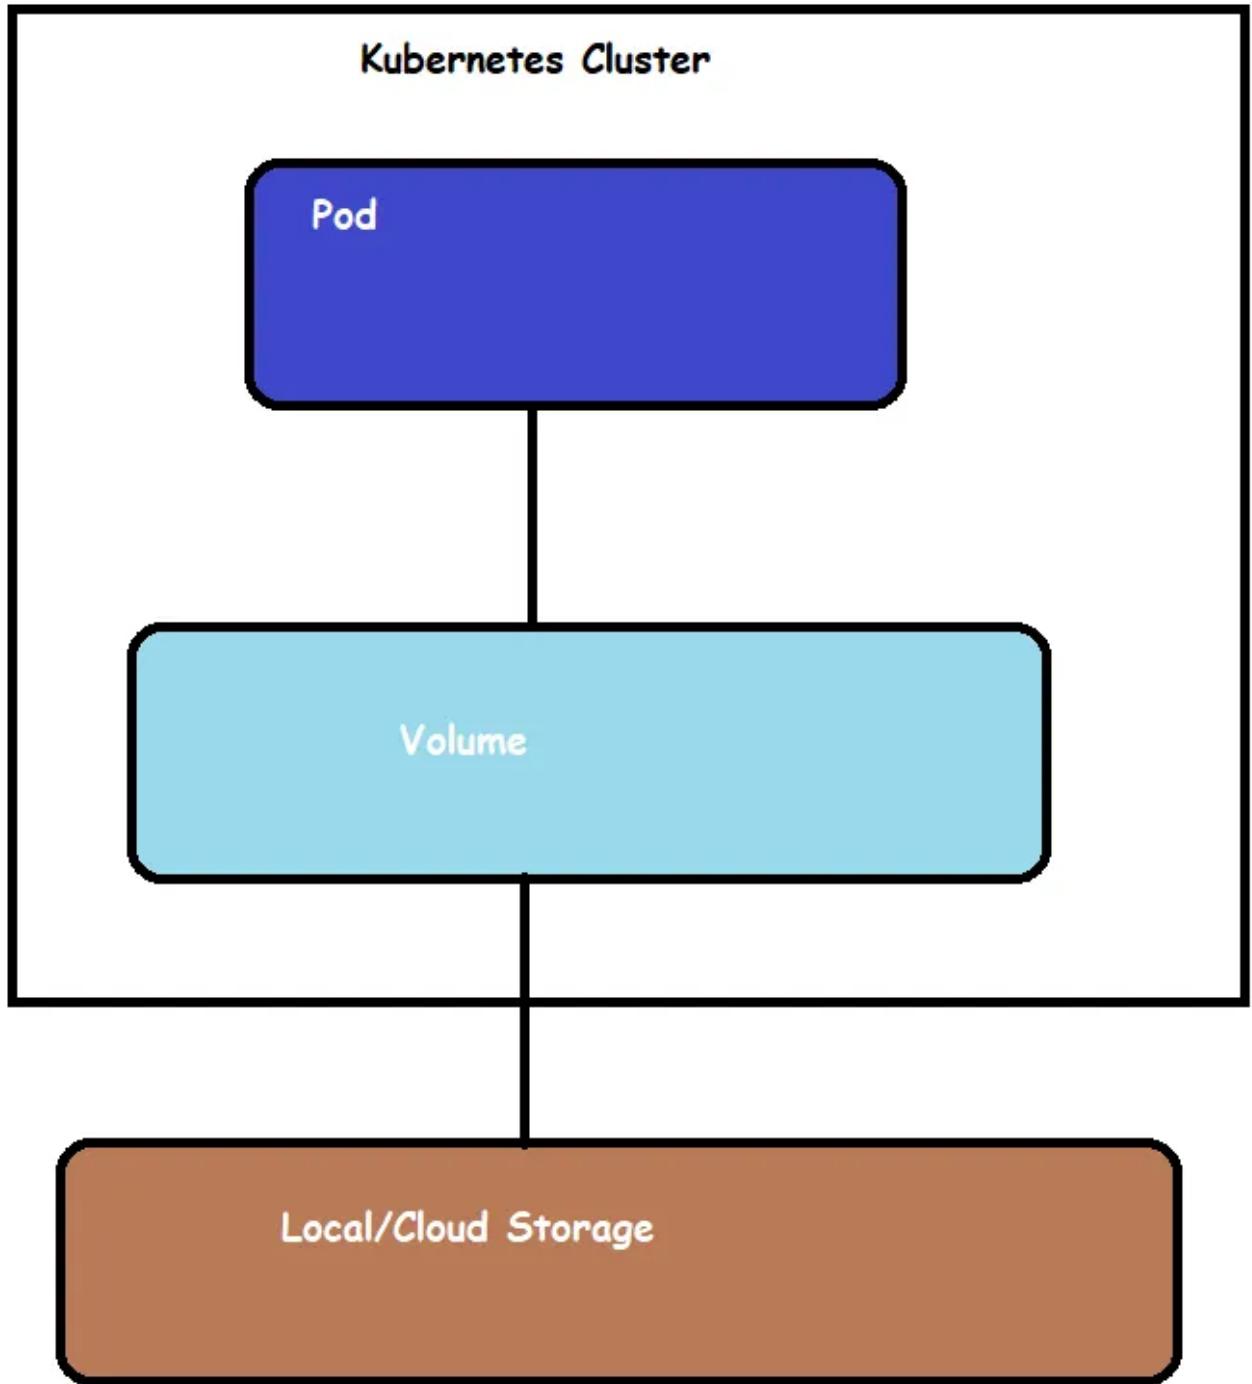
## Topic: Storage in K8s

- Both containers & Pods are ephemeral (when we delete the Pods/Containers data generated is deleted)
- Docker Volume separates storage from the life cycle of Container, So even if the container is deleted, data will still be available.
- Since Docker containers are created in the Pod, and the problem of storage has to be solved at Pod level in k8s as Pod provides kernel namespace to docker container.
- Pod has a container. The container in the pod stores the data and let's assume it is crashes, Pod will try to restart the container and in this case data will be lost. The new container will start with empty disk space allocated.



- We cannot rely on containers themselves even for temporary storage of data.

- K8s Volume abstraction solves the problem



- Volume is exposed to the applications which eventually store data on any type of storage (node/cloud).
- Lifetime of k8s volume is the same as the Pod's lifetime. As long as Pod is alive the containers within POD will get the same volume across restarts or new container Creation.
- How to use K8s Volume

- Volume is defined in Pod Spec
- Apply the k8s pod manifest with empty volume

```
---
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-emptydir
spec:
  containers:
    - image: qualitythought/gameoflife:07112020
      name: gol-container
      volumeMounts:
        - mountPath: /tmp
          name: tmp-volume
  volumes:
    - name: tmp-volume
      emptyDir: {}
```

```
qtkhaja@Azure:~$ kubectl apply -f podwithemptyvol.yaml
pod/pod-with-emptydir created
qtkhaja@Azure:~$
qtkhaja@Azure:~$ kubectl get po
NAME           READY   STATUS      RESTARTS   AGE
pod-with-emptydir  0/1   ContainerCreating   0          12s
qtkhaja@Azure:~$ kubectl get po -w
NAME           READY   STATUS      RESTARTS   AGE
pod-with-emptydir  1/1   Running   0          17s
^Cqtkhaja@Azure:~$ kubectl exec -it pod-with-emptydir -- /bin/bash
root@pod-with-emptydir:/usr/local/tomcat# df -h
Filesystem      Size  Used Avail Use% Mounted on
overlay         124G  21G  104G  17% /
tmpfs           64M    0   64M   0% /dev
tmpfs           3.5G   0   3.5G   0% /sys/fs/cgroup
/dev/sda1       124G  21G  104G  17% /tmp
shm              64M    0   64M   0% /dev/shm
tmpfs           3.5G  12K  3.5G   1% /run/secrets/kubernetes.io/serviceaccount
tmpfs           3.5G    0   3.5G   0% /proc/acpi
tmpfs           3.5G    0   3.5G   0% /proc/scsi
tmpfs           3.5G    0   3.5G   0% /sys/firmware
root@pod-with-emptydir:/usr/local/tomcat# 
```

```

ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/Volume$ ls
podwithemptyvol.yaml
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/Volume$ kubectl apply -f podwithemptyvol.yaml
pod/pod-with-emptydir created
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/Volume$ kubectl get pods
NAME           READY   STATUS    RESTARTS   AGE
pod-with-emptydir  0/1     ContainerCreating   0          8s
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/Volume$ kubectl get pods -w
NAME           READY   STATUS    RESTARTS   AGE
pod-with-emptydir  1/1     Running   0          14s
^Cubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/Volume$ kubectl exec -it pod-with-emptydir -- /bin/bash
root@pod-with-emptydir:/usr/local/tomcat# df -h
Filesystem      Size  Used Avail Use% Mounted on
overlay        80G  2.6G  78G  4% /
tmpfs          64M    0   64M  0% /dev
tmpfs          3.9G    0   3.9G  0% /sys/fs/cgroup
/dev/xvda1      80G  2.6G  78G  4% /tmp
shm            64M    0   64M  0% /dev/shm
tmpfs          3.9G  12K  3.9G  1% /run/secrets/kubernetes.io/serviceaccount
tmpfs          3.9G    0   3.9G  0% /proc/acpi
tmpfs          3.9G    0   3.9G  0% /proc/scsi
tmpfs          3.9G    0   3.9G  0% /sys/firmware
root@pod-with-emptydir:/usr/local/tomcat#

```

- [Refer Here](#) for official docs  
{<https://kubernetes.io/docs/concepts/storage/volumes/>}
- Lets use one more example of hostPath (Type of volume) [Refer Here](#)

```

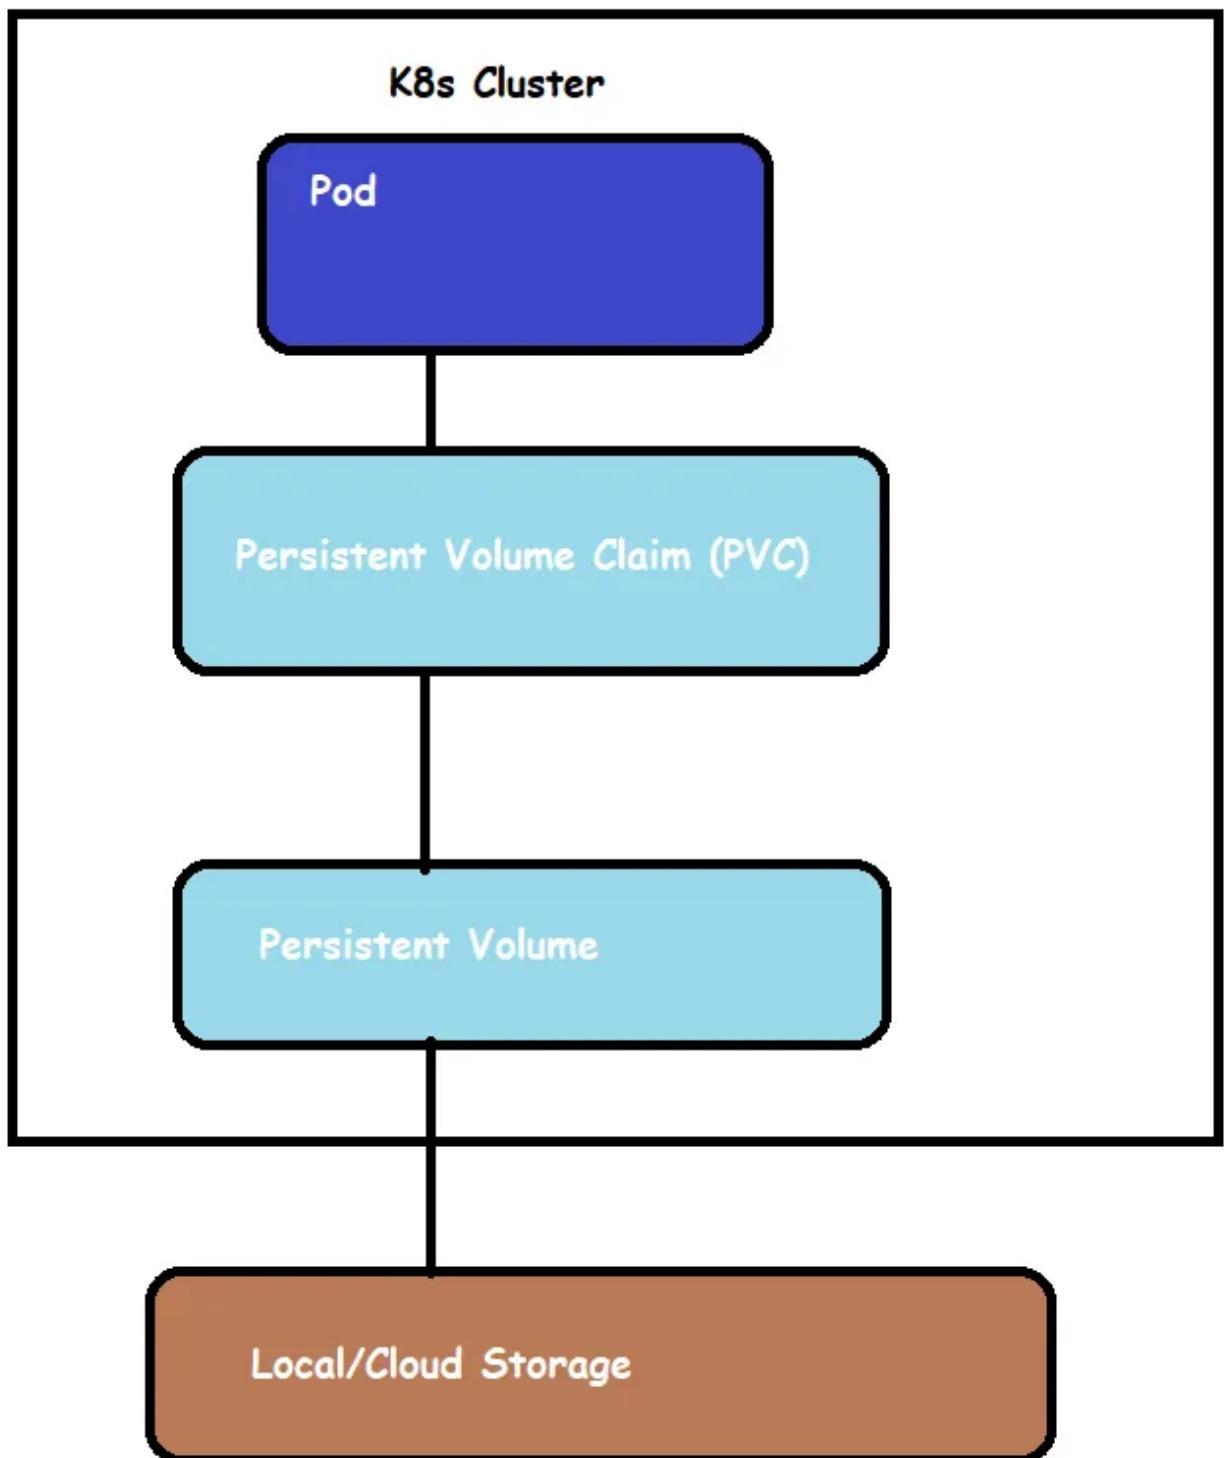
---
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-hostpath
spec:
  containers:
    - image: qualitythought/gameoflife:07112020
      name: gol-container
      volumeMounts:
        - mountPath: /tmp
          name: tmp-volume-host
  volumes:
    - name: tmp-volume-host
      hostPath:
        path: /tmp
        type: Directory

```

- Volumes such as emptyDir and hostPath get deleted when the pod using them is deleted or the pod is restarted. So we need to find a solution Where the data/storage lifetime is different than Pod's lifetime
- To Solve this Problem, K8s has Persistent Storage in the form of Persistent Volume (PV).

## Topic: Persistent Volumes & Persistent Volume Claims

- A PV is a K8s object that represents a block of storage in the cluster. This can be provisioned before hand by cluster admins or be dynamically provisioned.
- In Order to use PV, a Persistent Volume Claim (PVC) needs to be created. A PVC is a request for storage by user or by Pod. This request can be of specific size or specific access mode.



- Each PV belongs to a certain storage class. A Storage Class is a K8s Object that provides a way for administrators to describe different types of Storage.
- Volume mode has two Values *Block*, *FileSystem*
- Access Modes:
  - ReadWriteOnce (RWO): Mounted as read-write by a single node
  - ReadOnlyMany (ROX): Mounted as read-only by many nodes
  - ReadWriteMany (RWX): Mounted as Read-write by many nodes
- persistentVolume Reclaim Policy
  - **Retain** (Even if u delete your persistent volume the data will still be there )
  - **Recycle** (when we delete your persistent volume the store will be there but the data in the storage will be cleaned )
  - **Delete** (when ever we delete the persistence volume the storage/volume and data in that volume will be deleted)

## PVC in Azure

- Dynamic Provision Azure Disk [Refer Here](#) Azure Files [Refer Here](#)
- Static Provision Azure Disk [Refer Here](#) Azure Files [Refer Here](#)
- NFS in Azure [Refer Here](#)
- Azure Netapp [Refer Here](#)
- AWS [Refer Here](#)
- [Refer Here](#) for AWS Storage Classes Documentation

Storage Classes in Azure (In azure we have 4 different classes )

NAME	PROVISIONER	RECLAIMPOLICY	VOLUMEBINDINGMODE	ALLOWVOLUMEEXPANSION	AGE
azurefile	kubernetes.io/azure-file	Delete	Immediate	true	104m
azurefile-premium	kubernetes.io/azure-file	Delete	Immediate	true	104m
default (default)	kubernetes.io/azure-disk	Delete	Immediate	true	104m
managed-premium	kubernetes.io/azure-disk	Delete	Immediate	true	104m

## PVC in AWS

Storage Class in AWS (we have 1 storage class in aws)

NAME	PROVISIONER	RECLAIMPOLICY	VOLUMEBINDINGMODE	ALLOWVOLUMEEXPANSION	AGE
gp2 (default)	kubernetes.io/aws-ebs	Delete	WaitForFirstConsumer	false	100m
ubuntu@ip-172-31-10-23:~\$					

### Activity:

- Now let's create a 'my-sql' container, my-sql is a database so it require some kind of storage, so we will create a pod in that we will create a my-sql,
- Lets understand how to use mysql container in a Pod with Persistent Volume (in AWS and Azure)
- [Refer Here](#) to mysql Dockerfile which has volume at /var/lib/mysql

In azure:

- Lets use existing storage classes in Azure to create mysql-pod [Refer Here](#)

```
---  
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: mysql-azure-pvc  
spec:  
  accessModes:  
    - ReadWriteOnce  
  storageClassName: managed-premium  
resources:  
  requests:  
    storage: 1Gi  
---
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: mysql-azure  
spec:  
  containers:  
    - name: mysql-cont  
      image: mysql:8  
      volumeMounts:  
        - mountPath: /var/lib/mysql
```

```

name: mysql-volume
env:
  - name: MYSQL_ROOT_PASSWORD
    value: 'rootroot'
  - name: MYSQL_DATABASE
    value: 'openmrs'
  - name: MYSQL_USER
    value: qtdevops
  - name: MYSQL_PASSWORD
    value: qtdevops
ports:
  - name: mysqlport
    containerPort: 3306
    protocol: TCP

```

#### volumes:

- name: mysql-volume
 persistentVolumeClaim:
 claimName: mysql-azure-pvc

```

qtkhaja@Azure:~$ kubectl apply -f mysql-azure.yaml
persistentvolumeclaim/mysql-azure-pvc unchanged
pod/mysql-azure created
qtkhaja@Azure:~$ kubectl get pods -w
NAME           READY   STATUS      RESTARTS   AGE
mysql-azure     0/1     ContainerCreating   0          8s
pod-with-emptydir 1/1     Running     0          69m
^Cqtkhaja@Azure:~$ kubectl get pv
NAME               CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM               STORAGECLASS
REASON   AGE
pvc-42a3ff40-43c1-408d-8b89-c177f2e9683c  1Gi        RWO            Delete   Bound   default/mysql-azure-pvc   managed-premi
m       6m33s
qtkhaja@Azure:~$ 

```

Name	Type	Location
aks-agentpool-34786526-nsg	Network security group	East US
aks-agentpool-34786526-routetable	Route table	East US
aks-nodepool1-34786526-vms	Virtual machine scale set	East US
aks-vnet-34786526	Virtual network	East US
b3524f9d-e1d3-411d-b80a-b5e3827cf941	Public IP address	East US
kubernetes	Load balancer	East US
<b>kubernetes-dynamic-pvc-42a3ff40-43c1-408d-8b89-c177f2e9683c</b>	Disk	East US

## In AWS:

- for the manifest [Refer Here](#)

---

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pvc-aws
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: gp2
  resources:
    requests:
      storage: 1Gi
```

---

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-pvc-aws
spec:
  containers:
    - name: mysql-cont
      image: mysql:8
      volumeMounts:
        - mountPath: /var/lib/mysql
          name: mysql-volume
  env:
    - name: MYSQL_ROOT_PASSWORD
      value: 'rootroot'
    - name: MYSQL_DATABASE
      value: 'openmrs'
    - name: MYSQL_USER
      value: qtdevops
    - name: MYSQL_PASSWORD
      value: qtdevops
  ports:
```

```

- name: mysqlport
  containerPort: 3306
  protocol: TCP
volumes:
- name: mysql-volume
  persistentVolumeClaim:
    claimName: mysql-pvc-aws

```

```

ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/PersistentVolume$ kubectl apply -f mysql-aws.yaml
persistentvolumeclaim/mysql-aws-pvc created
pod/mysql-aws created
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/PersistentVolume$ kubectl get pv
NAME                                     CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM
STORAGECLASS     REASON   AGE
pvc-bbe2c3e5-e93a-4287-b1be-48c6ef0c8add  1Gi        RWO            Delete          Bound    default/mysql-aws-pvc
  gp2           1s
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/PersistentVolume$ kubectl get pods
NAME             READY   STATUS      RESTARTS   AGE
mysql-aws        0/1     ContainerCreating   0          14s
pod-with-emptydir 1/1     Running     0          70m
ubuntu@ip-172-31-10-23:~/KubernetesZone/Nov20/Storage/PersistentVolume$ 

```

Name	Volume ID	Size	Volume Type	IOPS	Snapshot	Created	Availability Zone	State
kubernetes-dynamic-pvc-bbe2c3e5-...	vol-0919eeef...	1 GiB	gp2	100		November 21, 2020...	us-west-2c	in-use
	vol-0009ffda...	80 GiB	gp2	240	snap-db5efac6...	November 21, 2020...	us-west-2c	in-use
eksctl	vol-0d2babbb...	8 GiB	gp2	100	snap-01c125a...	November 21, 2020...	us-west-2c	in-use

## K8s Volume Snapshots & Clones

- K8s Volumesnapshot represents the snapshot of volume on the storage system.
- Source can be specified as PVC [Refer Here](#)
- CSI Volume Cloning features add support for existing PVCs to clone the volume. [Refer Here](#)

## Topic: Replica set (K8s controller type-1)

- Replica Set will ensure that the number of replicas specified in the spec is maintained by k8s cluster.

**Replica Set**  
number of  
replicas 2



- K8s will try to maintain the desired state which is the number of replicas. If some pod gets deleted/crashed etc, k8s will create one more pod to maintain the desired state.

**Replica Set**  
number of  
replicas 2



- Pods might be allocated on different cluster nodes.
- So lets create a replica set with 2 replicas for the game of life. [Refer Here](#) for the changeset

---  
apiVersion: apps/v1

```

kind: ReplicaSet
metadata:
  name: gol-rs
spec:
  minReadySeconds: 5
  replicas: 2
  template:
    metadata:
      labels:
        app: gol
        version: "07112020"
    spec:
      containers:
        - image: qualitythought/gameoflife:07112020
          name: gol
          ports:
            - containerPort: 8080

```

- Now lets add selector for the labels



---

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: replicatio-set
spec:
  replicas: 2
  minReadySeconds: 10
  template:
    metadata:
      labels:
        app: gol
        version: "1.1"

```

spec:

  containers:

- image: qualitythought/gameoflife:07112020
- name: gol
- ports:
  - containerPort: 8080

  selector:

    matchLabels:

      app: gol

      version: "1.1"

- apply the replica set and explore
- kubectl get rs - to view the replica-set
- kubectl describe rs replication-set - to describe the replication set

```
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get rs
NAME           DESIRED   CURRENT   READY   AGE
hello-world-rest-api-55c678fc85   1         1         0       6d8h
replocatio-set   2         2         2       3m39s
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl describe rs replocatio-set
Name:           replocatio-set
Namespace:      default
Selector:       app=gol,version=1.1
Labels:          <none>
Annotations:    <none>
Replicas:       2 current / 2 desired
Pods Status:   2 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=gol
           version=1.1
Containers:
  gol:
    Image:      qualitythought/gameoflife:07112020
    Port:       8080/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:     <none>
    Volumes:    <none>
Events:
  Type      Reason            Age      From               Message
  ----      ----            --      --                --
  Normal   SuccessfulCreate  3m59s   replicaset-controller  Created pod: replocatio-set-d2xv6
  Normal   SuccessfulCreate  3m59s   replicaset-controller  Created pod: replocatio-set-ffpf6
root@ip-172-31-17-61:/mahi/k8s/rs#
```

```
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs -o wide
NAME   DESIRED   CURRENT   READY   AGE   CONTAINERS   IMAGES   SELECTOR
gol-rs  2         2         2       80s   gol          qualitythought/gameoflife:07112020   app=gol,version=07112020
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get pods -o wide
NAME        READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GATES
gol-rs-4px55  1/1    Running   0          108s  10.47.0.1   ip-172-31-17-40   <none>          <none>
gol-rs-mmxf5l 1/1    Running   0          108s  10.44.0.1   ip-172-31-27-107  <none>          <none>
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

Now lets shutdown one node and see the result

```

root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get rs -w
NAME      DESIRED   CURRENT   READY   AGE
hello-world-rest-api-55c678fc85  1         1         0       6d9h
replocatio-set          2         2         1       22m
replocatio-set          2         1         1       24m
replocatio-set          2         2         1       24m
replocatio-set          2         2         2       24m
replocatio-set          2         2         2       24m
^Croot@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods -o wide
NAME           READY   STATUS             RESTARTS   AGE   IP           NODE   NOMINATED NODE   READI
NESS GATES
hello-world-rest-api-55c678fc85-ndd85  0/1   ImagePullBackOff  0        37m  10.44.0.1  ip-172-31-12-116  <none>  <none>
>
replocatio-set-d2xv6          1/1   Terminating      0        26m  10.47.0.1  ip-172-31-15-137  <none>  <none>
>
replocatio-set-ffpfcc         1/1   Running          0        26m  10.44.0.2  ip-172-31-12-116  <none>  <none>
>
replocatio-set-z5cdk          1/1   Running          0        90s  10.44.0.3  ip-172-31-12-116  <none>  <none>
>

```

We can even scale the pods in replica set without modifying yaml file  
 kubectl scale replica set replication-set --replicas=5

NOTE:- it is not a recommended way

```

root@ip-172-31-17-61:/mahi/k8s/rs# kubectl scale replicaset replocatio-set --replicas=5
replicaset.apps/replocatio-set scaled
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get ods
error: the server doesn't have a resource type "ods"
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
hello-world-rest-api-55c678fc85  1         1         0       6d9h
replocatio-set          5         5         5       30m
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods
NAME           READY   STATUS             RESTARTS   AGE
hello-world-rest-api-55c678fc85-ndd85  0/1   ImagePullBackOff  0        42m
replocatio-set-d2xv6          1/1   Terminating      0        30m
replocatio-set-ffpfcc         1/1   Running          0        30m
replocatio-set-gmdlz          1/1   Running          0        18s
replocatio-set-pdr6q          1/1   Running          0        18s
replocatio-set-sjwrg          1/1   Running          0        18s
replocatio-set-z5cdk          1/1   Running          0        6m10s
root@ip-172-31-17-61:/mahi/k8s/rs# kubectl get pods -o wide
NAME           READY   STATUS             RESTARTS   AGE   IP           NODE   NOMINATED NODE   READI
NESS GATES
hello-world-rest-api-55c678fc85-ndd85  0/1   ImagePullBackOff  0        42m  10.44.0.1  ip-172-31-12-116  <none>
ne>
replocatio-set-d2xv6          1/1   Terminating      0        30m  10.47.0.1  ip-172-31-15-137  <none>
ne>
replocatio-set-ffpfcc         1/1   Running          0        30m  10.44.0.2  ip-172-31-12-116  <none>
ne>
replocatio-set-gmdlz          1/1   Running          0        26s  10.44.0.4  ip-172-31-12-116  <none>
ne>
replocatio-set-pdr6q          1/1   Running          0        26s  10.44.0.5  ip-172-31-12-116  <none>
ne>
replocatio-set-sjwrg          1/1   Running          0        26s  10.44.0.6  ip-172-31-12-116  <none>
ne>
replocatio-set-z5cdk          1/1   Running          0        6m18s 10.44.0.3  ip-172-31-12-116  <none>
ne>
root@ip-172-31-17-61:/mahi/k8s/rs# 

```

- when ever we need to scale the replicas is in two ways i) chaining the yaml file ii) using auto scaling imperative commands

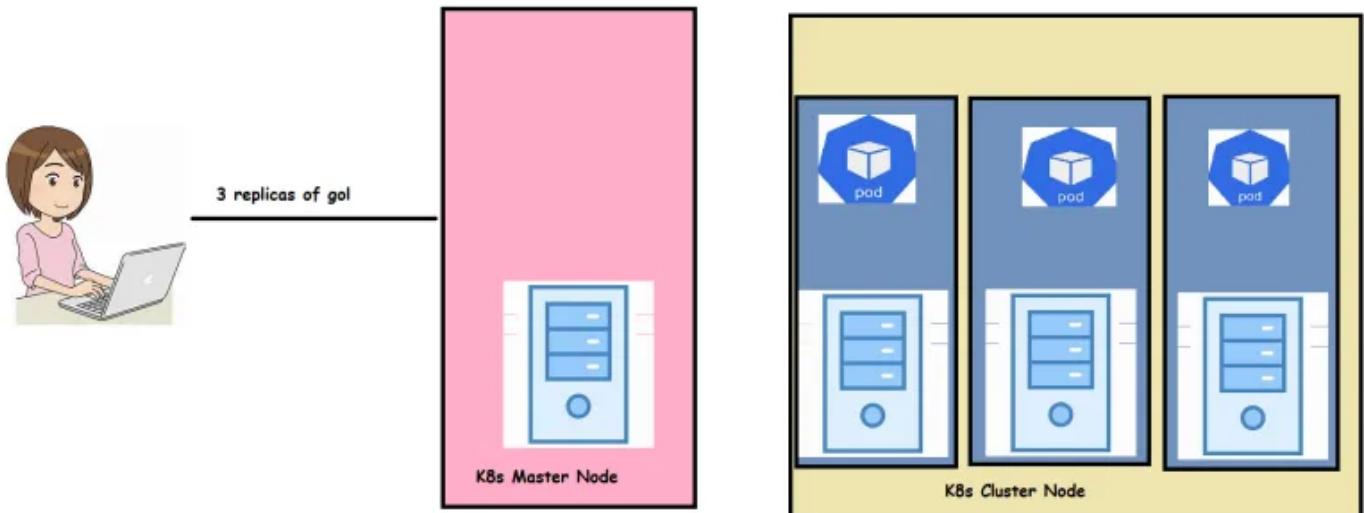
EX:- kubectl autoscale rs replocatio-set --min=2 --max=8 --cpu-percent=80

```

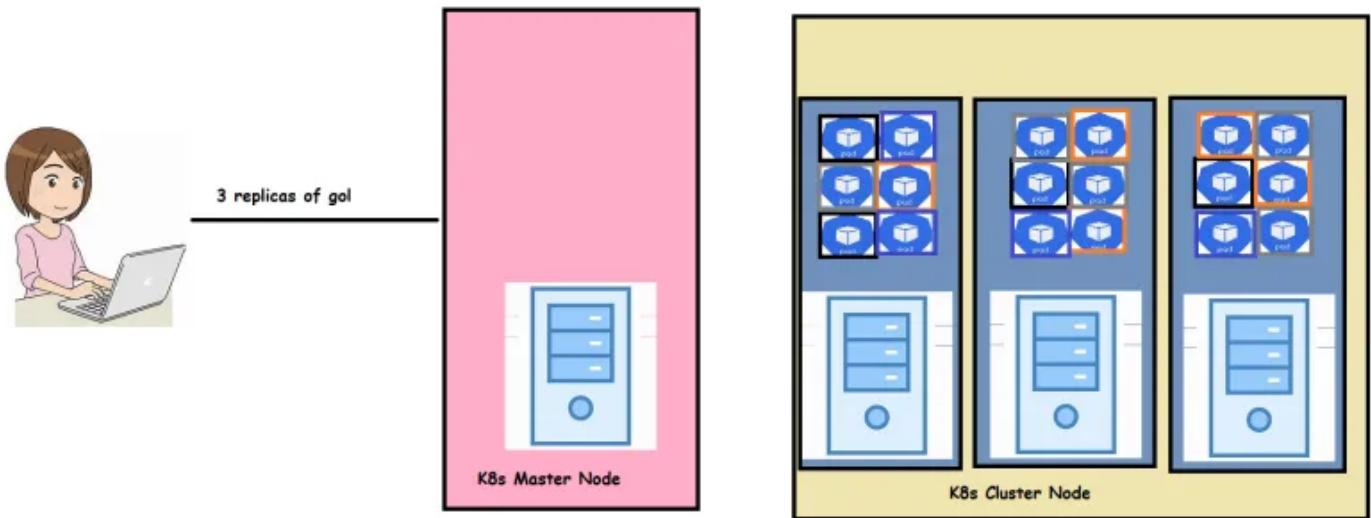
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl autoscale rs gol-rs --min=2 --max=5 --cpu-percent=80
horizontalpodautoscaler.autoscaling/gol-rs autoscaled
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
gol-rs    2          2          2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
gol-rs    2          2          2       11m
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get hpa
NAME      REFERENCE   TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
gol-rs   ReplicaSet/gol-rs <unknown>/80%   2          5          2          27s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

## Labels:-

- Basic workflow of creating replicas



- Giving any pod ip address to access the application is not good cuz any pod can be killed any time and any pod can create at any time.
- As k8s is used and application microservices run in the cluster, the number of pods running of nodes might grow in size and complexity.
- As shown in the below pic orange circled pods are belongs to one service for ex web server, Black circled pods are belongs to another service for ex app server..etc.,



- K8s has added foundational concepts of Labels and annotations to organization, mark and cross-index all of your resources to represent the groups.
- *Labels* are key/value pairs that can be attached to k8s objects such as pods and ReplicaSets. Labels are used by k8s cluster to group pods/other objects
- *Annotations* provide storage mechanism similar to labels which information that can be used by tools & libraries
- Labels provide identifying metadata



```
---
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: gol-rs-label
spec:
  minReadySeconds: 10
  replicas: 5
  selector:
```

```

matchLabels:
  app: gol
template:
metadata:
  labels:
    app: gol
    version: "07112020"
spec:
  containers:
    - image: qualitythought/gameoflife:07112020
      name: gol
      ports:
        - containerPort: 8080
          protocol: TCP

```

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl apply -f gol-rs-withlabels.yaml
replicaset.apps/gol-rs-label created
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
gol-rs-label   5         5         5       6s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
gol-rs-label-2hrnf  1/1     Running   0          13s
gol-rs-label-9gkjh  1/1     Running   0          13s
gol-rs-label-9l9pg  1/1     Running   0          13s
gol-rs-label-ln7p7  1/1     Running   0          13s
gol-rs-label-spxqj  1/1     Running   0          13s
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get pods -owide
NAME        READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED NODE   READINESS GAGE
ES
gol-rs-label-2hrnf  1/1     Running   0          20s  10.47.0.1   ip-172-31-17-40 <none>           <none>
gol-rs-label-9gkjh  1/1     Running   0          20s  10.44.0.1   ip-172-31-27-107 <none>           <none>
gol-rs-label-9l9pg  1/1     Running   0          20s  10.44.0.2   ip-172-31-27-107 <none>           <none>
gol-rs-label-ln7p7  1/1     Running   0          20s  10.47.0.2   ip-172-31-17-40 <none>           <none>
gol-rs-label-spxqj  1/1     Running   0          20s  10.47.0.3   ip-172-31-17-40 <none>           <none>
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

```

ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$ kubectl get pods --show-labels
NAME        READY   STATUS    RESTARTS   AGE   LABELS
gol-rs-label-2hrnf  1/1     Running   0          62s  app=gol,version=07112020
gol-rs-label-9gkjh  1/1     Running   0          62s  app=gol,version=07112020
gol-rs-label-9l9pg  1/1     Running   0          62s  app=gol,version=07112020
gol-rs-label-ln7p7  1/1     Running   0          62s  app=gol,version=07112020
gol-rs-label-spxqj  1/1     Running   0          62s  app=gol,version=07112020
ubuntu@ip-172-31-20-190:~/KubernetesZone/Nov20/replicaset$
```

- Pods can be selected by writing selectors

Some selector examples:

---

selector:

```

matchLabels:
  app: jenkins

```

```
---  
selector:  
  matchExpressions:  
    - {key: app, operator: In, values: ['jenkins', 'gol']}
```

```
---  
selector:  
  matchExpressions:  
    - {key: ver, operator: NotIn, values: ['1.0', '2.0']}
```

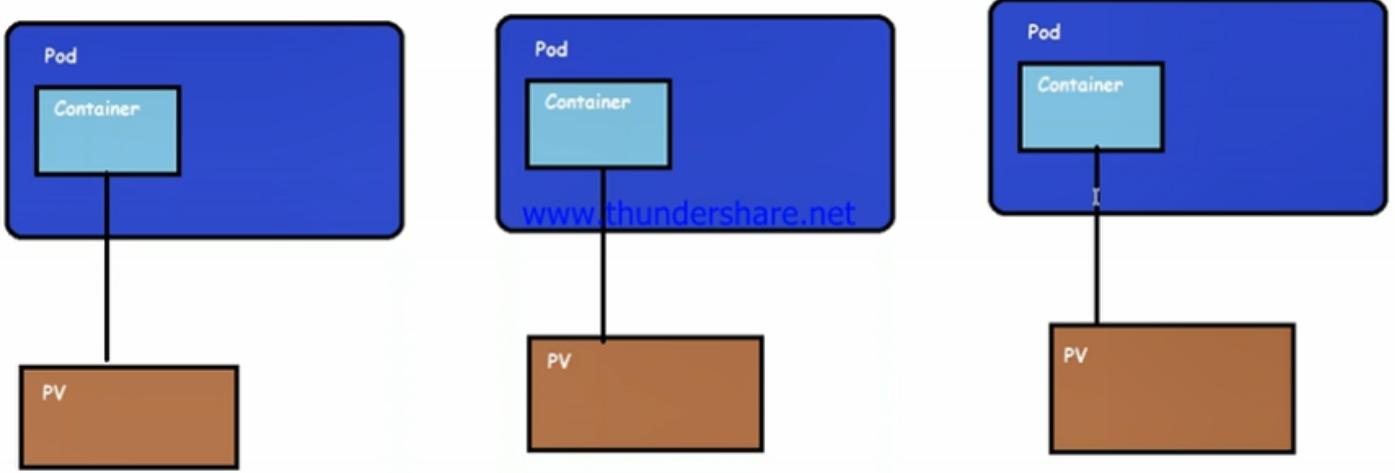
- If you want to all the api-resources in k8s cluster

## Service Discovery

- We are able to run our applications in pods
- We can dynamically scale pods by using autoscale command
- How do i access my application which is running in pods
- Service-discovery helps solving the problem of finding which processes are listening at which address for each service.
- We need to expose our services to the external world. As of now w.r.t networking we know each pod is having its own ip address.

## Topic: StatefulSet (K8s Controller type-2)

- Generally we map same storage disk/location to a application running in different pos/containers
- But if we each pod/container running for same application can maintain separate pvc/storage as shown in the below fig then we need stateful-set
- Stateful sets are used to stateful replicas



Stateful sets are used to stateful replicas [Refer Here](#)

```
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: stateful-demo
spec:
  replicas: 3
  serviceName: jenkins-svc
  selector:
    matchLabels:
      app: jenkins
  template:
    metadata:
      labels:
        app: jenkins
        version: "2.60.3"
    spec:
      containers:
        - image: jenkins:2.60.3
          name: jenkins
        ports:
          - containerPort: 8080
            protocol: TCP
```

```

qtkhaja@Azure:~$ kubectl apply -f Statefulset.yaml
error: error validating "Statefulset.yaml": error validating data: ValidationError(StatefulSet.spec): missing required field "serviceName" in io.k8s.api.apps.v1.StatefulSetSpec; if you choose to ignore these errors, turn validation off with --validate=false
qtkhaja@Azure:~$ kubectl apply -f Statefulset.yaml
statefulset.apps/stateful-demo created
qtkhaja@Azure:~$ kubectl api-resources | grep State*
qtkhaja@Azure:~$ kubectl api-resources | grep State
statefulsets          sts          apps          true      StatefulSet
healthstates          health       azmon.container.insights   true      HealthState
qtkhaja@Azure:~$ kubectl get sts
NAME        READY   AGE
stateful-demo 1/3    31s
qtkhaja@Azure:~$ kubectl get sts -w
NAME        READY   AGE
stateful-demo 3/3    36s

```

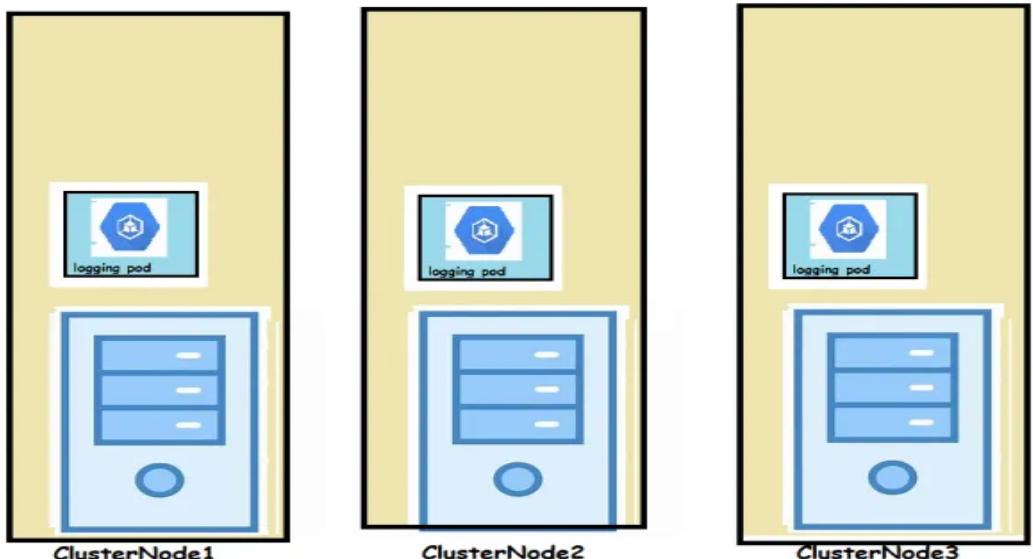
```

root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get sts
NAME        READY   AGE
statefulset-demo 3/3    3m3s
root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
statefulset-demo-0 1/1    Running   0          2m31s
statefulset-demo-1 1/1    Running   0          2m29s
statefulset-demo-2 1/1    Running   0          2m27s
root@ip-172-31-27-144:/mahi/k8s/statefull# kubectl get pods -o wide
NAME           READY   STATUS    RESTARTS   AGE   IP                                     NODE
statefulset-demo-0 1/1    Running   0          2m45s  192.168.49.114   ip-192-168-34-80.ec2.internal <none>
statefulset-demo-1 1/1    Running   0          2m43s  192.168.21.0    ip-192-168-8-224.ec2.internal <none>
statefulset-demo-2 1/1    Running   0          2m41s  192.168.42.229  ip-192-168-34-80.ec2.internal <none>
root@ip-172-31-27-144:/mahi/k8s/statefull#

```

## Topic: DaemonSet (K8s Controller type-3)

- Daemonsets are used to manage creation of a Particular Pod on all or selected nodes in k8s cluster
- Use Cases:
  - Logging agent
  - Monitoring agent



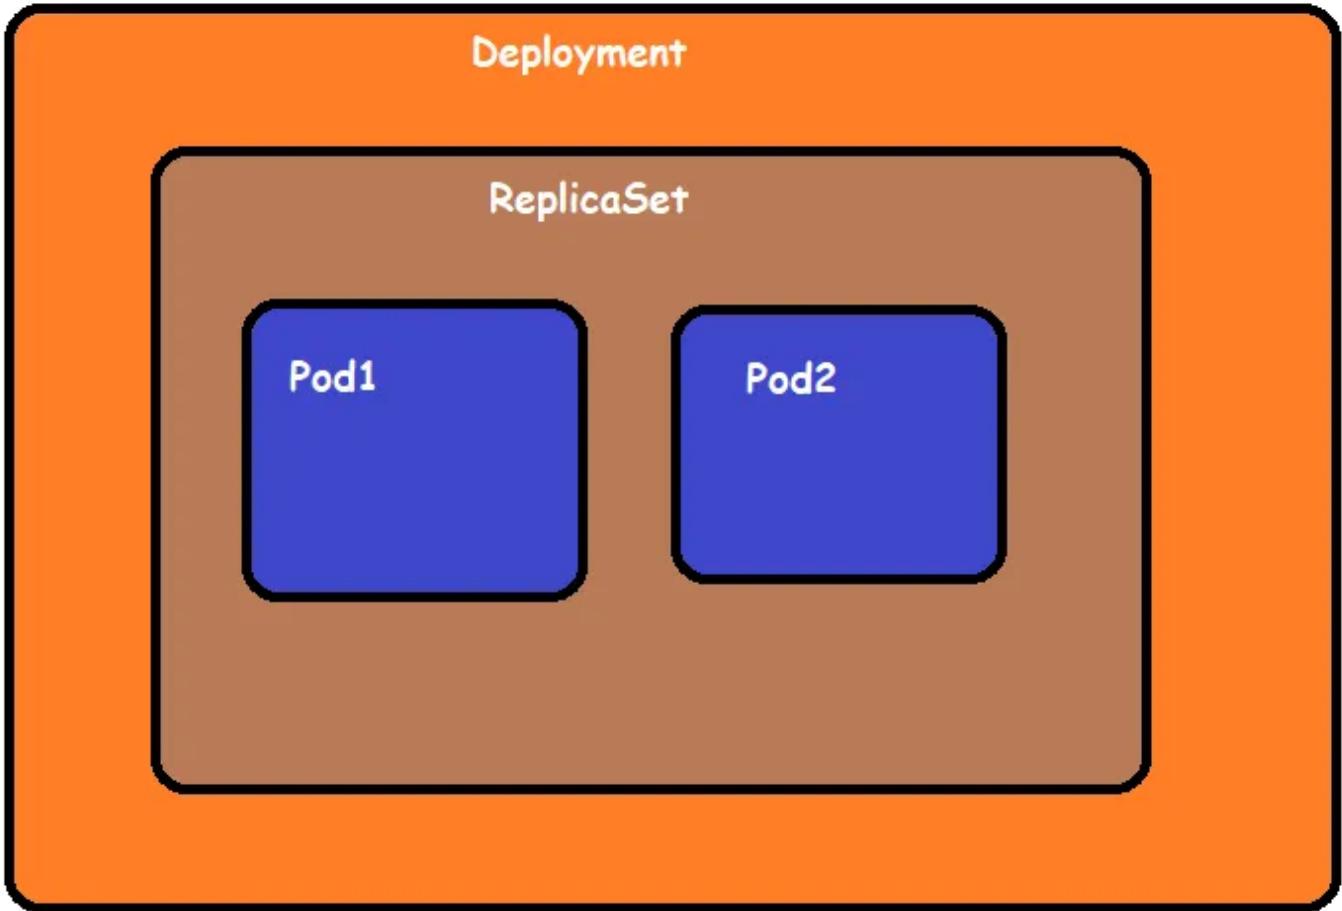
Apply the spec [Refer Here](#)

```
---  
apiVersion: apps/v1  
kind: DaemonSet  
metadata:  
  name: daemonset-demo  
spec:  
  selector:  
    matchLabels:  
      app: alpine  
  template:  
    metadata:  
      labels:  
        app: alpine  
    spec:  
      containers:  
        - name: alpine-cont  
          image: alpine  
          args:  
            - /bin/sh  
            - -C  
            - sleep 1d
```

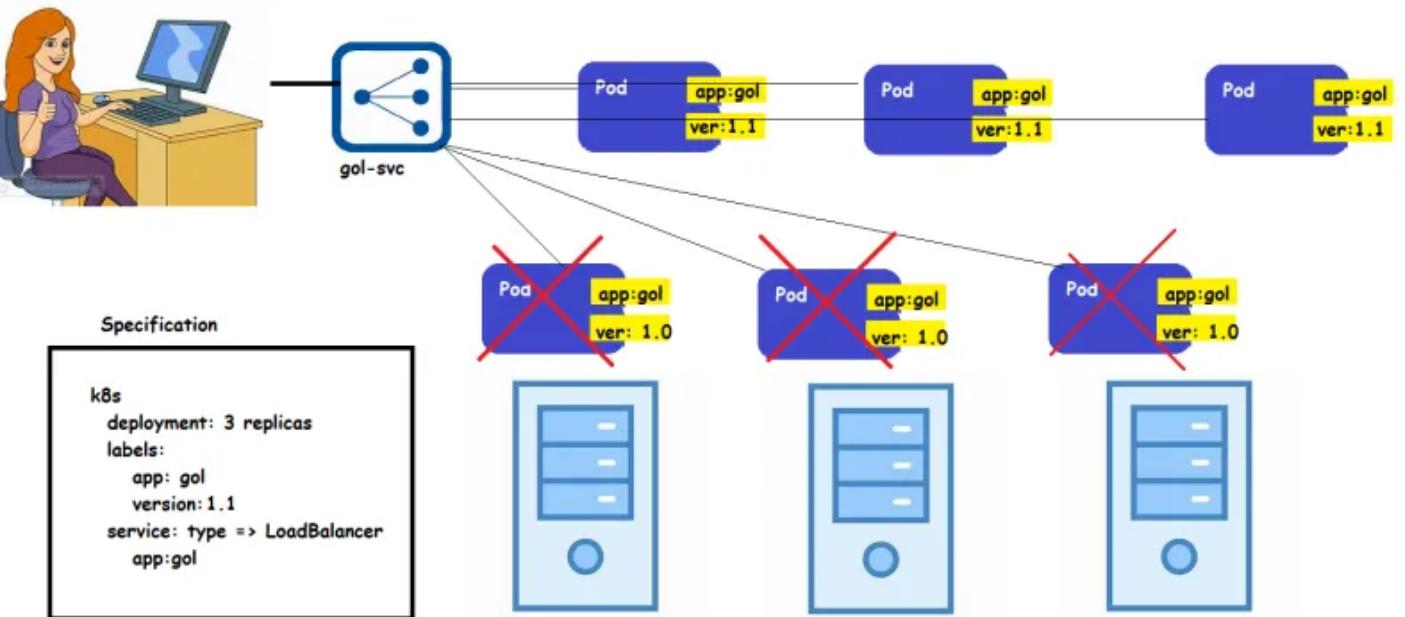
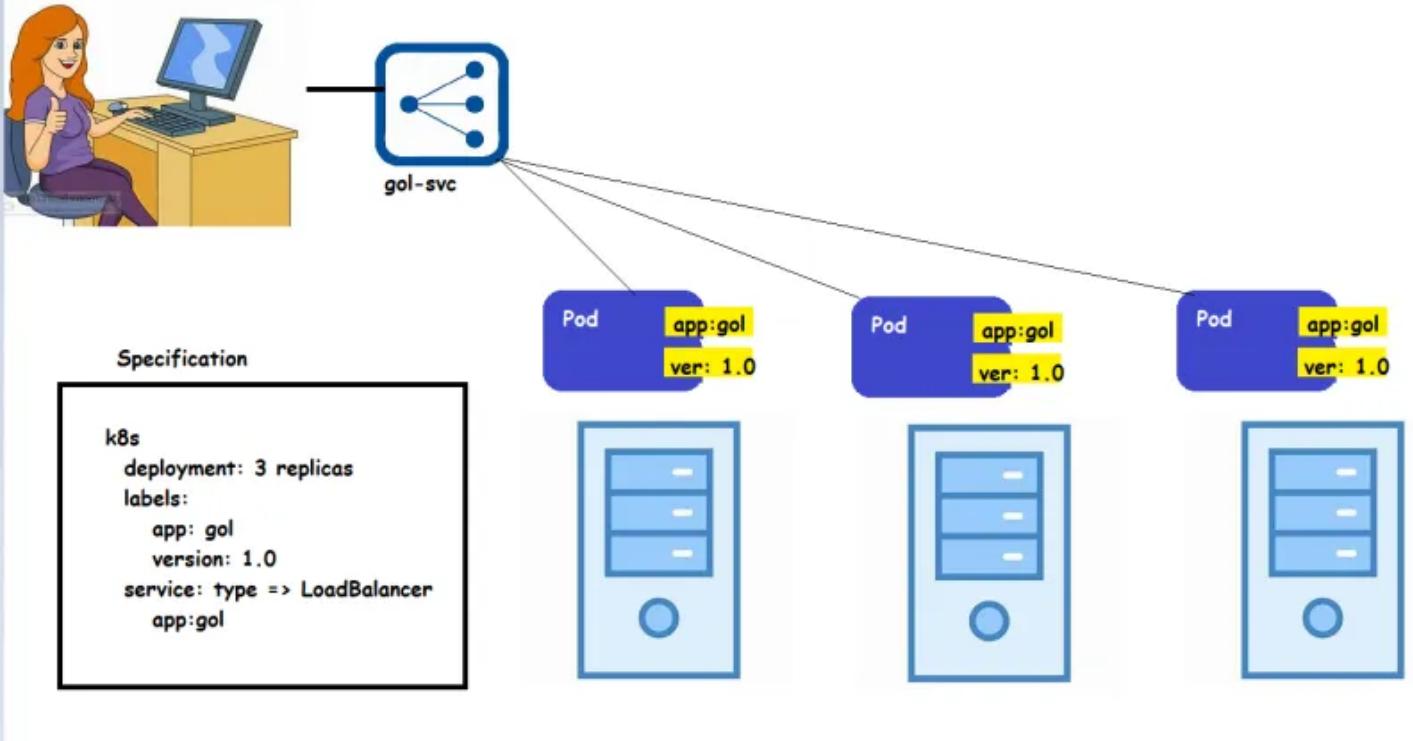
```
qtkhaja@Azure:~$ kubectl apply -f daemonset.yaml  
daemonset.apps/daemonset-demo created  
qtkhaja@Azure:~$ kubectl api-resources | grep daemon  
daemonsets          ds           apps           true       DaemonSet  
qtkhaja@Azure:~$ kubectl get ds  
NAME      DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR   AGE  
daemonset-demo  3         3         3        3           3           <none>        32s  
qtkhaja@Azure:~$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
daemonset-demo-4vmbn  1/1     Running   0          39s  
daemonset-demo-b66cp  1/1     Running   0          39s  
daemonset-demo-rwsz9  1/1     Running   0          39s  
qtkhaja@Azure:~$ kubectl get pods -o wide  
NAME          READY   STATUS    RESTARTS   AGE   IP           NODE          NOMINATED NODE   READINESS  
GATES  
daemonset-demo-4vmbn  1/1     Running   0          46s  10.244.1.3  aks-nodepool1-34786526-vmss000001  <none>        <none>  
daemonset-demo-b66cp  1/1     Running   0          46s  10.244.2.3  aks-nodepool1-34786526-vmss000002  <none>        <none>  
daemonset-demo-rwsz9  1/1     Running   0          46s  10.244.0.17 aks-nodepool1-34786526-vmss000000  <none>        <none>  
qtkhaja@Azure:~$ [ ]
```

## Topic: K8S Deployments (K8s Controller type-4)

Deployment is a k8s object that acts as a wrapper around ReplicaSet and makes it easier to use



- The major reason for Deployment is
  - it maintains the history of revisions (i.e. history of your application Deployments)
  - it makes it easy to roll back to previous revisions
  - it gives a way to do zero downtime deployments
- Deployment Strategy:
  - Rolling Update:
    - This give an option to deploy your application without having downtime
    - This gives maxUnavailable, maxSurge
  - Recreate:
    - This option kills all the running pods and recreates new pods



- [Refer Here](#) for deployment commands
- Now lets write a sample deployment [Refer Here](#)

```
---
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: gameoflife-deploy
  labels:
    app: gameoflife
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: gameoflife
      version: "1.0"
  template:
    metadata:
      labels:
        app: gameoflife
        version: "1.0"
    spec:
      containers:
        - image: qualitythought/gameoflife:07112020
          name: gol-dev
          ports:
            - containerPort: 8080
              protocol: TCP
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: gol-dev
spec:
  type: LoadBalancer
  selector:
    app: gameoflife
  ports:
    - targetPort: 8080
      port: 80
```

```

qtkhaja@Azure:~$ kubectl apply -f gol-deployment.yaml
deployment.apps/gameoflife-deploy created
qtkhaja@Azure:~$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   1/3     3           1          11s
qtkhaja@Azure:~$ kubectl get deploy -w
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   2/3     3           2          16s
gameoflife-deploy   3/3     3           3          16s
^Cqtkhaja@Azure:~$ kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
gameoflife-deploy-754b466d7f-hlk9b   1/1     Running   0          33s
gameoflife-deploy-754b466d7f-krflj    1/1     Running   0          33s
gameoflife-deploy-754b466d7f-xkvw1   1/1     Running   0          33s
qtkhaja@Azure:~$ kubectl apply -f gol-svc.yaml
service/gol-dev created
qtkhaja@Azure:~$ kubectl get svc -w
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
gol-dev   LoadBalancer   10.0.83.222   <pending>      80:31136/TCP   6s
kubernetes   ClusterIP      10.0.0.1      <none>        443/TCP      15h
gol-dev   LoadBalancer   10.0.83.222   52.151.219.235   80:31136/TCP   14s

```

```

qtkhaja@Azure:~$ kubectl describe rs gameoflife-deploy
Name:           gameoflife-deploy-754b466d7f
Namespace:      default
Selector:       app=gameoflife,pod-template-hash=754b466d7f,version=1.0
Labels:         app=gameoflife
               pod-template-hash=754b466d7f
               version=1.0
Annotations:    deployment.kubernetes.io/desired-replicas: 3
                 deployment.kubernetes.io/max-replicas: 4
                 deployment.kubernetes.io/revision: 1
Controlled By: Deployment/gameoflife-deploy
Replicas:       3 current / 3 desired
Pods Status:    3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=gameoflife
           pod-template-hash=754b466d7f
           version=1.0

```

```

qtkhaja@Azure:~$ kubectl rollout history deployment gameoflife-deploy
deployment.apps/gameoflife-deploy
REVISION  CHANGE-CAUSE
1          <none>
qtkhaja@Azure:~$ 

```

```

qtkhaja@Azure:~$ kubectl apply -f gol-deployment.yaml --record
deployment.apps/gameoflife-deploy created
qtkhaja@Azure:~$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   3/3     3           3          8s
qtkhaja@Azure:~$ kubectl apply -f gol-deployment.yaml --record
deployment.apps/gameoflife-deploy configured
qtkhaja@Azure:~$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   2/3     2           2          81s
qtkhaja@Azure:~$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   2/3     2           2         103s
qtkhaja@Azure:~$ kubectl get deploy -w
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   2/3     3           2         106s
gameoflife-deploy   3/3     3           3         106s
gameoflife-deploy   3/3     3           3         106s
gameoflife-deploy   2/3     3           2         106s
gameoflife-deploy   3/3     3           3         108s

```

```

qtkhaja@Azure:~$ kubectl rollout history deployment gameoflife-deploy
deployment.apps/gameoflife-deploy
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=gol-deployment.yaml --record=true
2          kubectl apply --filename=gol-deployment.yaml --record=true

qtkhaja@Azure:~$ kubectl rollout undo deployment gameoflife-deploy --to-revision=1
deployment.apps/gameoflife-deploy rolled back
qtkhaja@Azure:~$ kubectl rollout status
error: required resource not specified
qtkhaja@Azure:~$ kubectl rollout status deployment gameoflife-deploy
deployment "gameoflife-deploy" successfully rolled out
qtkhaja@Azure:~$ kubectl get deploy
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
gameoflife-deploy   3/3     3           3         7m55s
qtkhaja@Azure:~$ 

```

- Follow the complete application creation in k8s [Refer Here](#)

# Topic:- K8s Jobs and CronJobs

- Let's assume you want to run a Pod which runs a task and then terminates gracefully, To solve these kind of Problems k8s has Jobs & CronJobs
- The difference b/w k8s job & CronJob is CronJobs can be scheduled to run on a particular time frame
- Run the job as mentioned in the changeset [Refer Here](#)

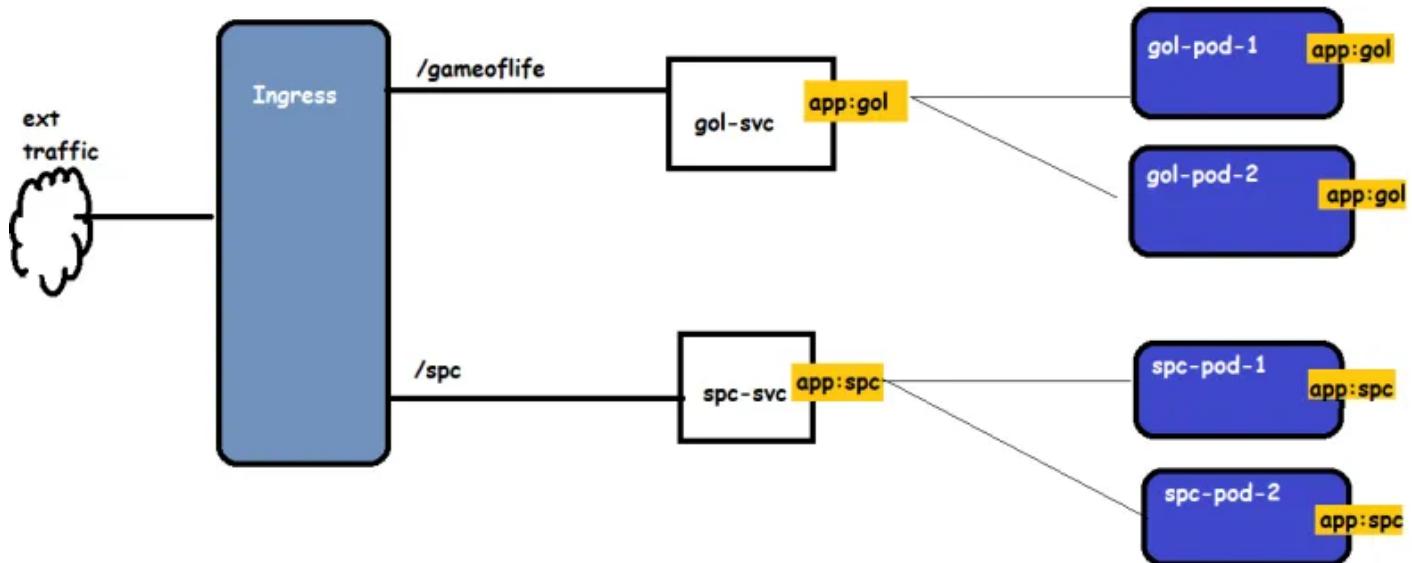
---

```
apiVersion: batch/v1
kind: Job
metadata:
  name: one-time-job
spec:
  template:
    spec:
      restartPolicy: OnFailure
      containers:
        - name: ubuntu-container
          image: ubuntu
          args:
            - /bin/bash
            - --c
            - sleep 10m
```

```
qtkhaja@Azure:~$ kubectl apply -f jobdemo.yaml
job.batch/one-time-job created
qtkhaja@Azure:~$ kubectl get jobs
NAME      COMPLETIONS   DURATION   AGE
one-time-job  0/1        4s         4s
qtkhaja@Azure:~$ kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
gameoflife-deploy-754b466d7f-wkq6b  1/1     Running   0          8h
gameoflife-deploy-754b466d7f-x8c5k  1/1     Running   0          8h
gameoflife-deploy-754b466d7f-zklnw  1/1     Running   0          8h
one-time-job-kwddj                1/1     Running   0          19s
qtkhaja@Azure:~$
```

## Topic: K8s Ingress

- Ingress is an object in k8s that is used manage external access to services in the k8s cluster. Ingress provides layer 7 load balancing, so it will be aware of url's
- Refer Here {<https://docs.microsoft.com/en-us/azure/aks/ingress-basic>}

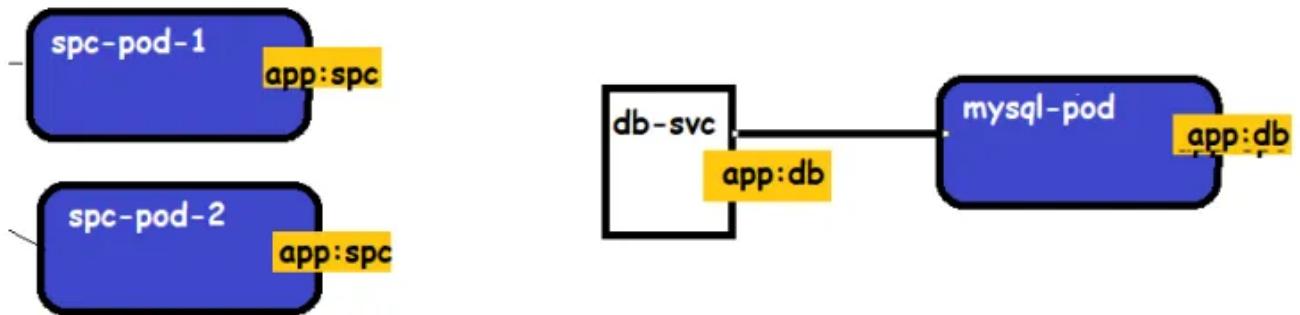


### Scenario:

- I have a application component which is running in a pod
- I have a database pod
- Application Pod is configured to scale automatically



- Now application pod needs to access db
- Approach 1:
  - Create a service for db



- Whenever you add a service an environmental variables will be added in all pods containing svc ip and port

```
DIRSTACK=()
EUID=0
GOL_DEV_PORT=tcp://10.0.83.222:80
GOL_DEV_PORT_80_TCP=tcp://10.0.83.222:80
GOL_DEV_PORT_80_TCP_ADDR=10.0.83.222
GOL_DEV_PORT_80_TCP_PORT=80
GOL_DEV_PORT_80_TCP_PROTO=tcp
GOL_DEV_SERVICE_HOST=10.0.83.222
GOL_DEV_SERVICE_PORT=80
GPG_KEYS='05AB33110949707C93A279E3D3EFE6B686867BA6 07E48665A34DCFAE522
1FBE7D8F78B25E055DDEE13C370389288584E7 61B832AC2F1C5A90F0F9B00A1C506407
B966EBA586F72C284D731FABEE A27677289986DB50844682F8ACB77FC2E86E29AC A90
B21E8933C60243 F3A04C595DB5B6A5F1ECA43E3B7BBB100D811BBE F7DA48BB64BCB84
GROUPS=()
HISTFILE=/root/.bash_history
```

## Approach 2:

- Use a Config Map. *Config Map* is a k8s object that allows us to define the application-related data.
- Config maps are key value pairs
- You can create config maps from
  - literal value

- file
- directory
- Lets create a sample config map

```
kubectl create configmap demovalue-map --from-literal=db-ip=10.0.83.222
```

---

```
apiVersion: v1
kind: Pod
metadata:
  name: configmapdemo
spec:
  containers:
    - name: configmapdemo
      image: alpine
      command: ["/bin/bash", "-c", "sleep 1d"]
      envFrom:
        - configMapRef:
            name: demovalue-map
```

```
qtkhaja@Azure:~$ kubectl apply -f configmapinpod.yaml
pod/configmap-pod created
qtkhaja@Azure:~$ kubectl get pods
NAME                      READY   STATUS    RESTARTS   AGE
configmap-pod           1/1     Running   0          5s
gameoflife-deploy-754b466d7f-wkq6b  1/1     Running   0          8h
gameoflife-deploy-754b466d7f-x8c5k  1/1     Running   0          8h
gameoflife-deploy-754b466d7f-zklnw  1/1     Running   0          8h
one-time-job-kwddj            0/1     Completed  0          30m
qtkhaja@Azure:~$ █
```

```
qtkhaja@Azure:~$ kubectl exec -it configmap-pod -- /bin/sh
/ # set
GOL_DEV_PORT='tcp://10.0.83.222:80'
GOL_DEV_PORT_80_TCP='tcp://10.0.83.222:80'
GOL_DEV_PORT_80_TCP_ADDR='10.0.83.222'
GOL_DEV_PORT_80_TCP_PORT='80'
GOL_DEV_PORT_80_TCP_PROTO='tcp'
GOL_DEV_SERVICE_HOST='10.0.83.222'
GOL_DEV_SERVICE_PORT='80'
HISTFILE='/root/.ash_history'
HOME='/root'
HOSTNAME='configmap-pod'
IFS='
```

NOTE:- In K8s we also have secret which are much like config maps but they can store secret of size 1kb in base 64 encoded form

## Topic: Helm

- Helm is a Package manager for k8s
- Helm charts represent a k8s package
- Installing helm [Refer Here](#)
- Lets execute an example using helm in AKS [Refer Here](#)
- In k8s we write yaml manifests, helm simplifies and makes manifests reusable
- [Refer Here](#)
- Journey of App from docker to k8s [Refer Here](#)